

Introduction to Computer Architecture

Chapter 2

Instructions: Language of the Computer - 1

Hyungmin Cho

Department of Computer Science and Engineering
Sungkyunkwan University

Levels of Program Code

■ High-level code

- ❖ Closer to the problem domain
- ❖ Better productivity & portability

■ Assembly code

- ❖ Human-readable representation of the machine code
- ❖ Can use symbolic names (labels)

■ Machine code

- ❖ Instructions encoded in binary format
- ❖ No symbols

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for RISC-V)

```
swap:
    slli x6, x11, 3
    add  x6, x10, x6
    ld   x5, 0(x6)
    ld   x7, 8(x6)
    sd   x7, 0(x6)
    sd   x5, 8(x6)
    jalr x0, 0(x1)
```

Assembler

Binary machine
language
program
(for RISC-V)

```
00000000001101011001001100010011
00000000011001010000001100110011
00000000000000110011001010000011
00000000100000110011001110000011
00000000011100110011000000100011
00000000010100110011010000100011
0000000000000001000000001100111
```

Instruction Set

- The collection of instructions of a computer
 - ❖ Arithmetic operations (**add**, **sub**, ...)
 - ❖ Logic operations (**and**, **or**, **xor**, ...)
 - ❖ Multiplication / Division
 - ❖ Load from memory / Store to memory
 - ❖ Jump / branch

 - ❖ Exception handling
 - ❖ Floating point

Chapter 2

- We will learn the RISC-V ISA
 - ❖ Understand basic instructions in RISC-V ISA
 - ❖ Writing simple programs in RISC-V assembly language
 - ❖ Binary representations of the instructions (machine code)

The RISC-V Instruction Set

- Used as the textbook's primary example for explaining the CPU architecture
- Developed by UC Berkeley as an **open ISA**
- Currently managed by RISC-V International (riscv.org)
- RISC-V simulators
 - ❖ Spike
 - <https://github.com/riscv-software-src/riscv-isa-sim>
 - ❖ RARS Simulator
 - <https://github.com/TheThirdOne/rars>
 - ❖ Complete list of simulators supporting RISC-V ISA:
 - <https://riscv.org/exchange/software/>

Arithmetic Operations

- Add and subtract operations have three operands
 - ❖ One destination and two sources

add a, b, c # $a \leftarrow b + c$

- All arithmetic operations follow this form

Arithmetic Operation Example

- Suppose you have “add” and “sub” instructions with two operands.

- Example C code:

```
f = (g + h) - (i + j); // 4 operands
```

- Corresponding RISC-V assembly code:


```
add t0, g, h    # t0 = g + h    // temporary
add t1, i, j    # t1 = i + j    // temporary
sub f, t0, t1   # f = t0 - t1
```

Registers

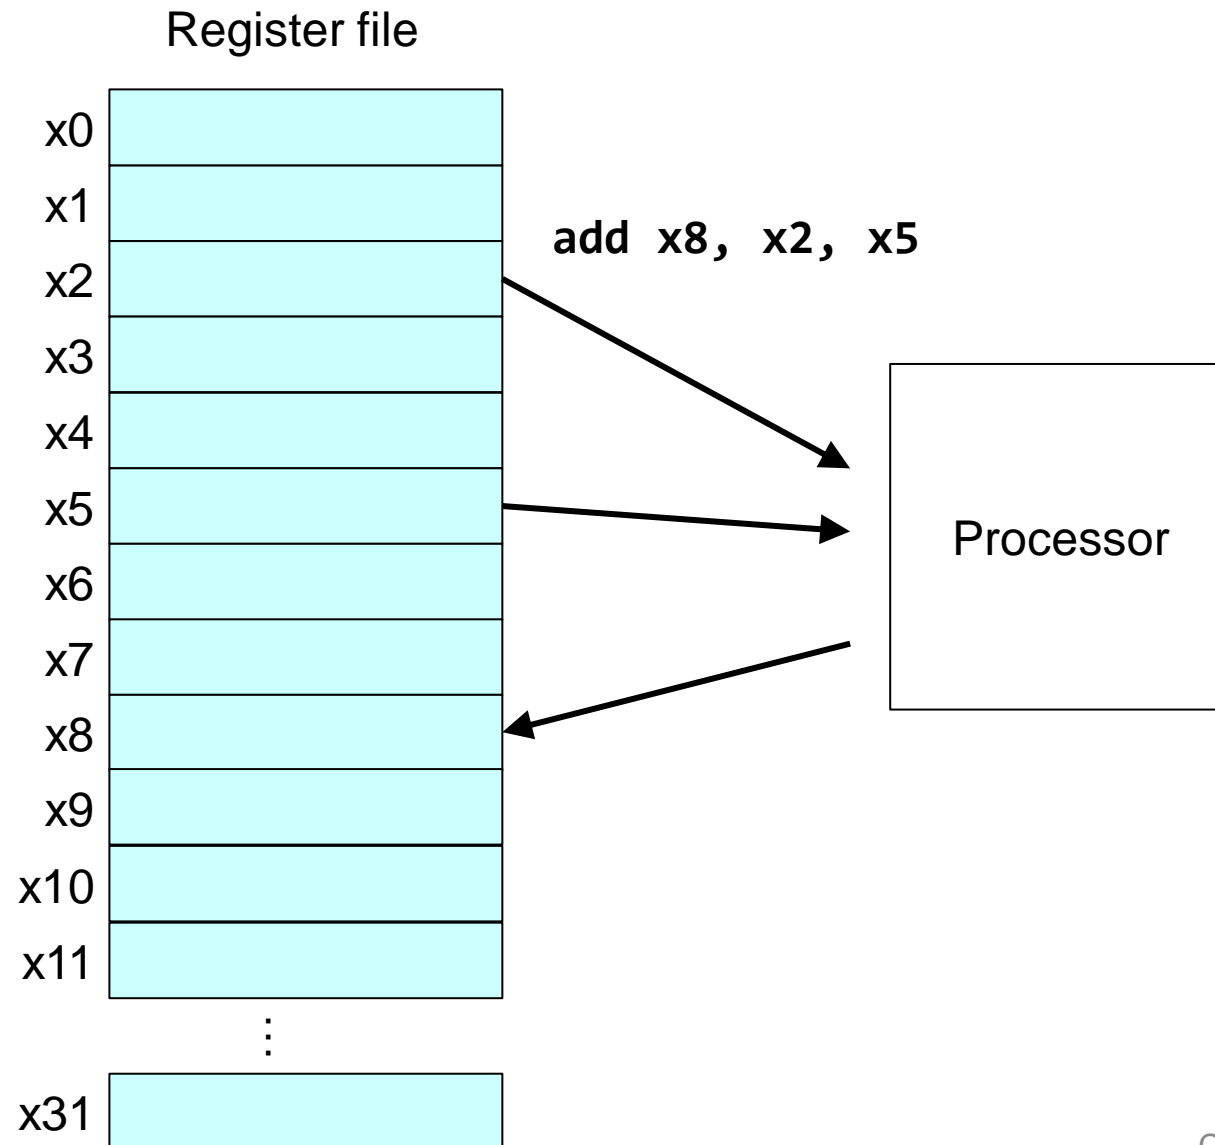
- Small, high-speed storage inside the CPU
- Registers hold data during the execution
 - ❖ Arithmetic instructions use register operands
- RISC-V (**RV32I**) ISA has a 32 registers, and the size of each register is 32bits.
 - ❖ The collection of those 32 registers is called “register file”
 - ❖ Each register has its own number (name): Numbered from x0 to x31
 - ❖ A 32-bit value is referred to as a “word”
 - ❖ Please note that the textbook is based on the 64-bit version (RV64I)
 - ❖ This lecture will be using the **32-bit RV32I**, not the 64-bit version

Registers

- RISC-V (RV32I) : 32 x 32-bit registers



RISC-V Registers



Memory

- Can't fit the entire data in registers
- Memory is the main storage for program data
- To apply arithmetic operations on data in memory...
 - ❖ Load: memory → registers
 - ❖ Compute: (source) register → (destination) register
 - ❖ Store: register → memory
- Memory is byte **addressed**
 - ❖ Each address: 1 byte (8-bit)

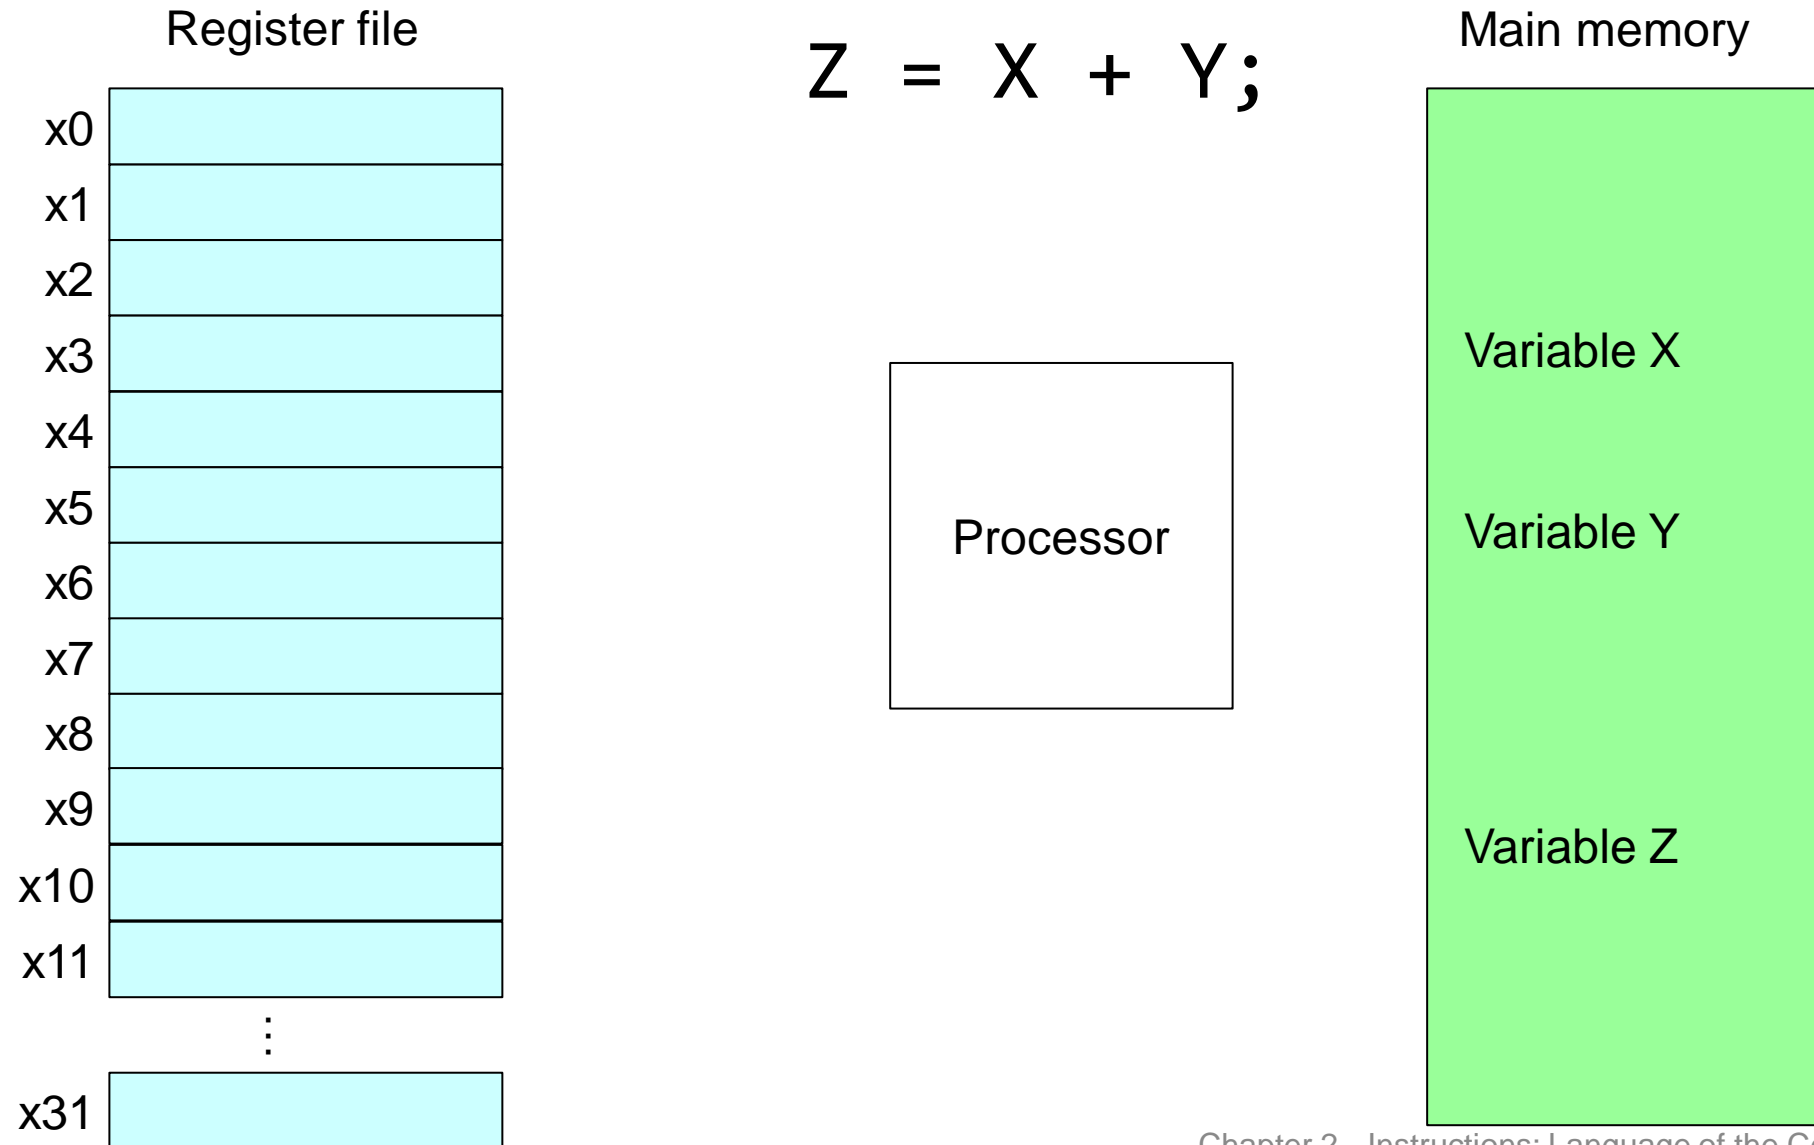
Register vs. Memory

- Registers are much faster than memory
- Can't perform computations directly on the data within memory
 - ❖ Need Load → Compute → Store steps
 - ❖ If need multiple operations, Load → Compute 1 → Compute 2 → ... → Store
 - ❖ Note: while some CISC-style CPUs support memory operands, such instructions will be broken down into multiple sub-operations internally.

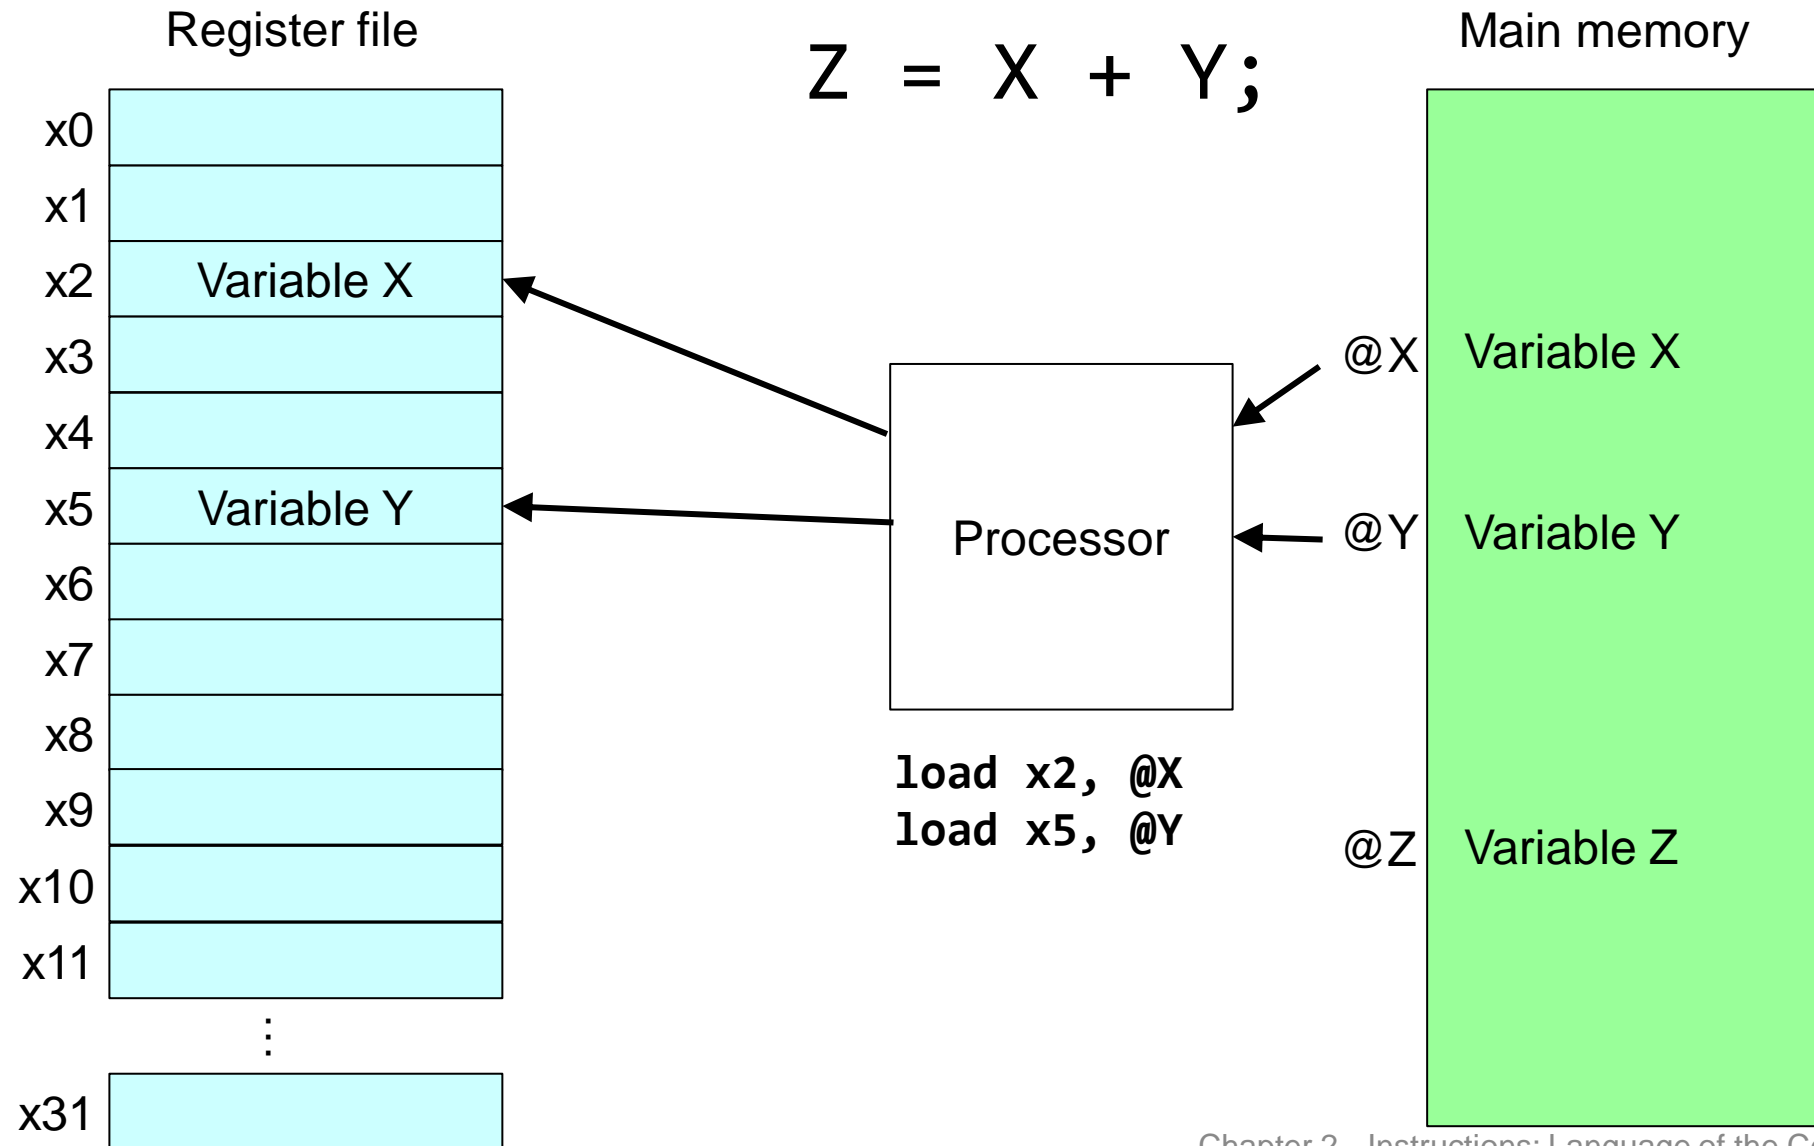
Register Optimization: Crucial for Performance.

- A good compiler should minimize memory access.
 - ❖ If an operation can be performed entirely using the registers only, the compiler can avoid memory accesses.
 - ❖ If the data size is larger than the registers, the compiler should prioritize frequently accessed data for registers and store less frequently accessed data in memory.

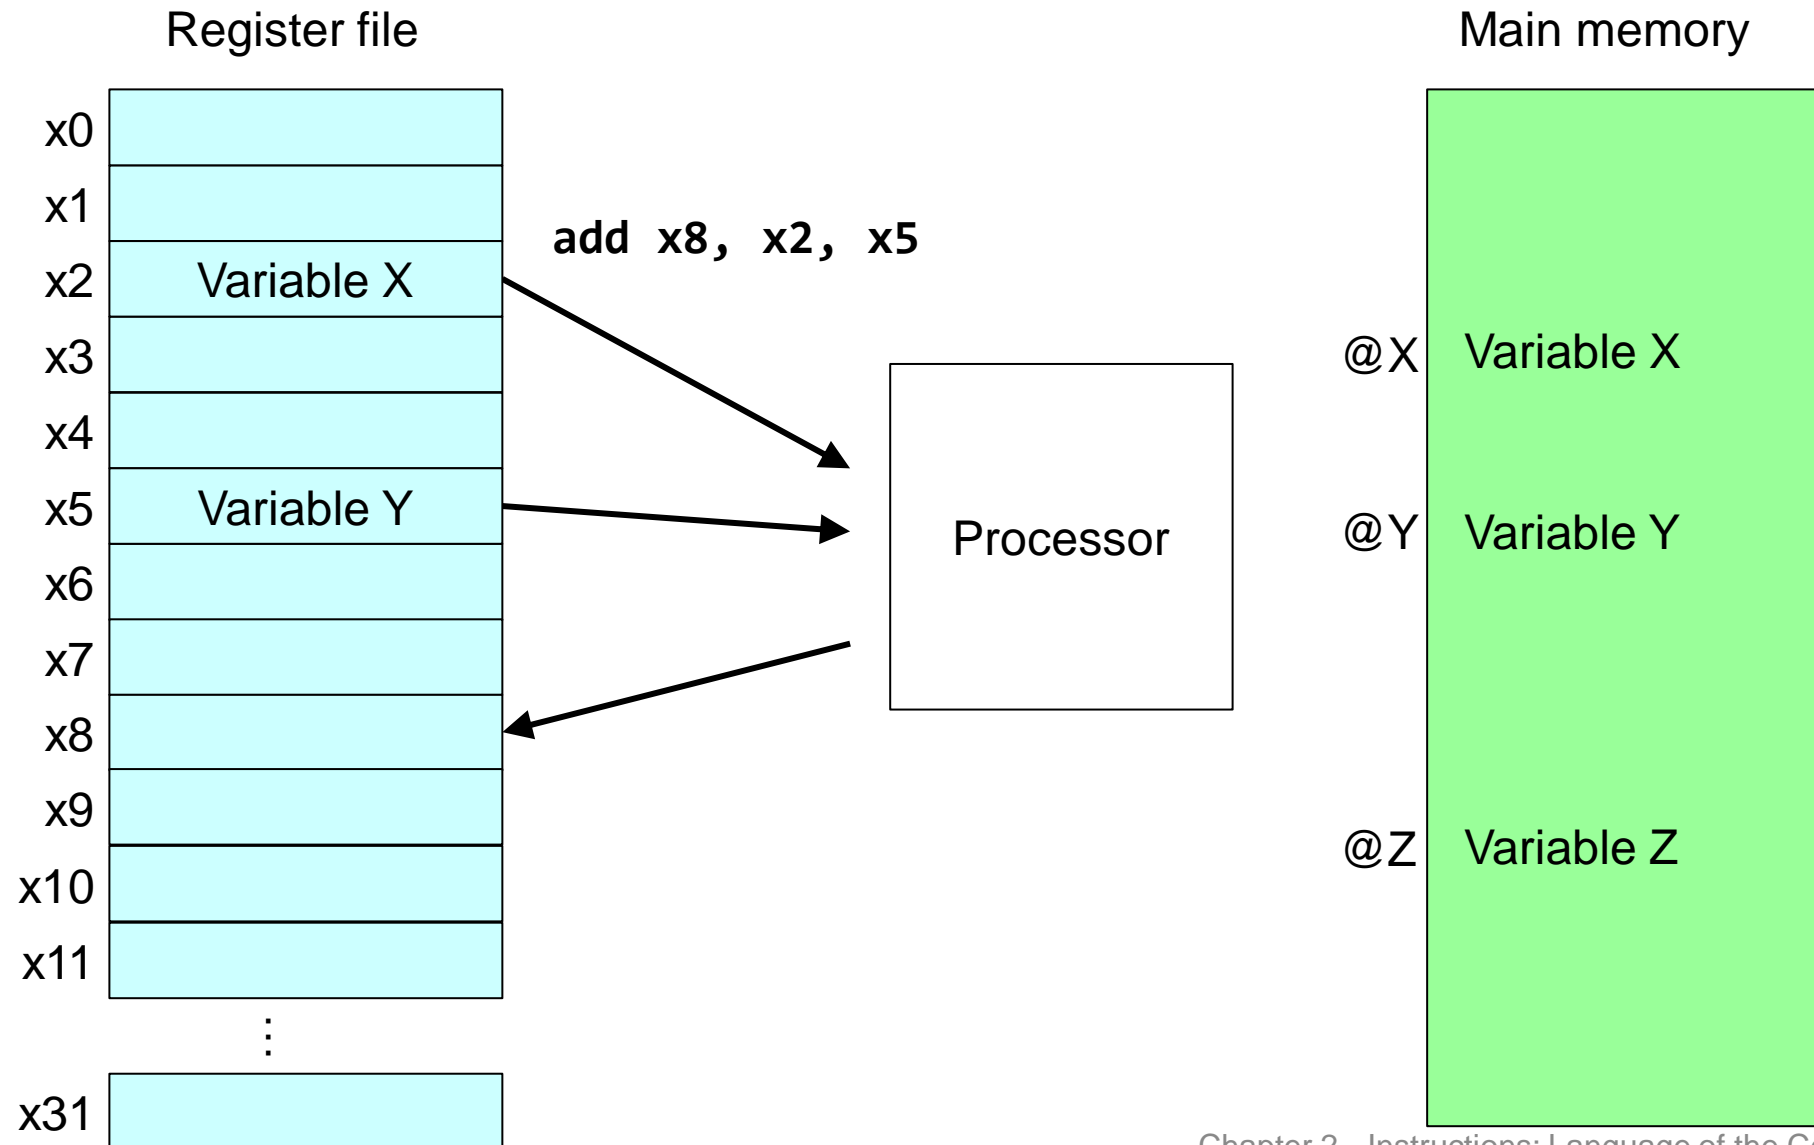
Using Memory Values



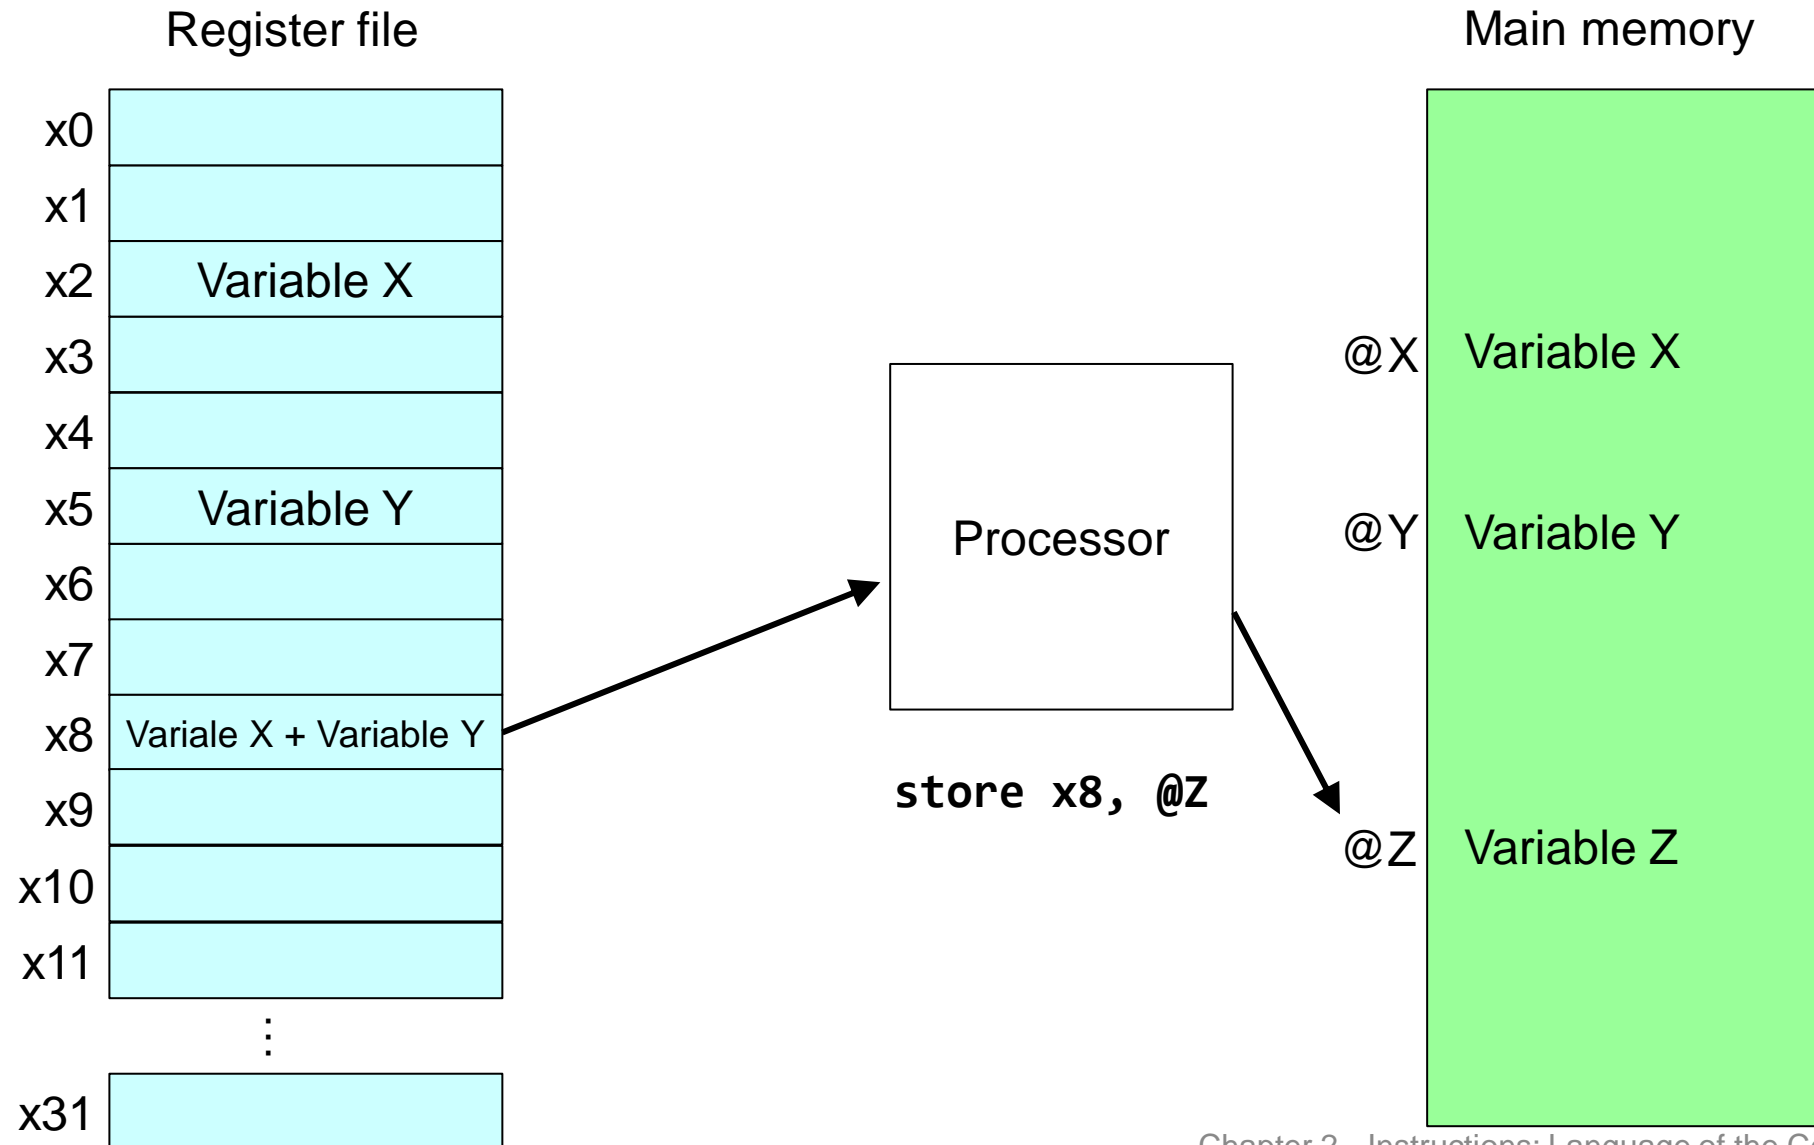
Using Memory Values



Using Memory Values



Using Memory Values



RISC-V Register Usage

Register Number	Name	Usage
x0	zero	Constant 0 (hardwired)
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7, x28-x31	t0 - t6	Temporaries
x8	s0 / fp	Frame pointer
x9, x18-x27	s1 - s11	Saved registers
x10-x11	a0 - a1	Function arguments / results
x12-x17	a2 - a7	Function arguments

Register Operand Example

- Example C code:

`f = (g + h) - (i + j);`

❖ Suppose `f, g, h, i, j` in `x19, x20, x21, x22, x23`

- Corresponding RISC-V assembly code:

`add x5, x20, x21`

`add x6, x22, x23`

`sub x19, x5, x6`

RISC-V Arithmetic Instructions

- RISC-V arithmetic instructions: `add`, `sub`, `addi`

Sample Code

- Suppose the registers have the following values

x1: 100, x2: 300, x3: 500, x4: 700

- What would be the value of a4(x14) after executing the following code?

add x11, x1, x2	//x11: 400
sub x12, x0, x3	//x12: -500
add x13, x11, x12	//x13: -100
add x14, x13, x4	//x14: 600

The Number of Registers is Finite

- What happens if you have no more registers?
 - ❖ You will eventually run out of registers if you use a new register for each instr.

- Reuse registers that are no longer used

```
add x1, x2, x3
sub x5, x0, x4
add x6, x1, x5
add x8, x6, x7
```

```
add x1, x2, x3
sub x5, x0, x4
add x1, x1, x5
add x8, x1, x7
```

- If all registers will be used later? Save some register values in memory
 - ❖ “Register spill”
 - ❖ Need to avoid as much as possible

Immediate Operands

- **Constant** data included in an instruction

addi x22, x23, 4

- No subtract immediate instruction (~~subi~~)

❖ Instead, just use a negative integer value.

addi x22, x23, -1

The Constant Zero

- RISC-V register 0 (x0, or 'zero') is the constant 0
 - ❖ Cannot be overwritten
- Useful for common operations
 - ❖ e.g., move (copy) between registers
add x1, x2, zero
 - ❖ Loading a constant (immediate) value to a register
addi x1, x0, 10

Load Immediate Value

- **li** instruction

`li xd, immediate_value (32-bit)`

- Example:

`li x1, 100`

`li x2, 0x10`

`add x3, x1, x2`

- Not a real instruction (*pseudo-instruction*)

Sample Code (2)

- x1: 100, x2: 300, x3: 500, x4: 700

```
li x1, 100
li x2, 300
li x3, 500
li x4, 0x2BC          //0x2bc = 700
add x10, x1, x2
sub x11, x0, x3
add x12, x10, x11
add x13, x12, x4
```

Multiply (1)

- Example: $11_2 \times 11_2 = 110_2 + 11_2 = 1001_2$
- 32-bit \times 32-bit \rightarrow Result can be up to 64-bit
- Need 2 destination registers to completely hold the result!

Multiply (2)


- “**M**” Extension for Multiply / Division
 - ❖ The base “RV32I” RISC-V ISA does not contain multiply / division instructions
 - ❖ “RV32I**M**” version supports multiply / division instructions
- `mul x3, x1, x2`
 - ❖ Multiply x1 and x2, and put the lower 32 bits to x3
- `mulh x3, x1, x2`
 - ❖ Multiply x1 and x2, and put the upper 32 bits to x3

Multiply (3)

- `mulh x3, x1, x2`
 - ❖ Multiply `x1` and `x2`, and put the upper 32 bits to `x3`
 - ❖ `x1` and `x2` contains signed numbers
- `mulhu x3, x1, x2`
 - ❖ Multiply `x1` and `x2`, and put the upper 32 bits to `x3`
 - ❖ `x1` and `x2` contains unsigned numbers
- `mulhsu x3, x1, x2`
 - ❖ Multiply `x1` and `x2`, and put the upper 32 bits to `x3`
 - ❖ `x1`: signed
 - ❖ `x2`: unsigned

Unsigned Binary Integers

- Given an n -bit number



$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$


- Example:

$$\begin{aligned} & 0000\ 0000\ 0000\ \dots\ 0000\ 1011_2 \\ &= 0 \times 2^{31} + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 \\ &= 0 + \dots + 8 + 0 + 2 + 1 = 11_{10} \end{aligned}$$

- Using 32 bits: 0 to +4,294,967,295

Signed Binary Integers



- How to represent a negative integer?





- Sign Magnitude: Any downside?

2's-Complement Signed Integers


- Given an $n-1$ bit positive integer... to make it a negative number with the same absolute value,



Step 1: negate all bits



Step 2: add 1



2's-Complement Signed Integers

- Given an n-bit number

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$

- Example (32-bit):

$$\begin{aligned} & 1111 \ 1111 \ \dots \ 1111 \ 1100_2 \\ &= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0 \\ &= -2,147,483,648 + 2,147,483,644 = -4_{10} \end{aligned}$$

- Using 32 bits: $-2,147,483,648$ to $+2,147,483,647$

2's-Complement Advantages

- Representing 1 more number (-2^{n-1})
- Only 1 representation of zero (0)
 - ❖ 0000 0000 0000
- Use the same adder logic with the unsigned integer values

Singed / Unsigned Multiplication

■ $0xFFFF_FFFF \times 0x0000_0002$

❖ If signed numbers: $-1 \times 2 = -2$

➤ -2 in 64 bit: FFFF FFFF FFFF FFFE

❖ If unsigned numbers: $4,294,967,295 \times 2 = 8,589,934,590$

➤ 8,589,934,590 in 64 bit: 0000 0001 FFFF FFFE

Division

- **div x3, x1, x2**

- ❖ Divide x1 by x2, and put the quotient to x3

- **rem x3, x1, x2**

- ❖ Divide x1 by x2, and put the remainder to x3

- **divu / remu**

- ❖ Unsigned versions

Logical Operations

- Instructions for bitwise manipulation

Operation	C	Java	RISC-V
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise XOR	^	^	xor, xori
Bitwise NOT	~	~	
Shift left	<<	<<	slli
Shift right	>>	>>>	srli

AND Operations

- Useful to mask bits in a word
 - ❖ Select some bits, clear others to 0

and x9, x10, x11

x10	0000	0000	0000	0000	0000	1101	1100	0000
x11	0000	0000	0000	0000	0011	1100	0000	0000
x9	0000	0000	0000	0000	0000	1100	0000	0000

OR Operations

- Useful to include bits in a word
 - ❖ Set some bits to 1, leave others unchanged

or **x9**, **x10**, **x11**

x10	0000	0000	0000	0000	0000	1101	1100	0000
x11	0000	0000	0000	0000	0011	1100	0000	0000
x9	0000	0000	0000	0000	0011	1101	1100	0000

XOR Operations

- Differencing operation

- ❖ Set some bits to 1, leave others unchanged

```
xor x9, x10, x11 // if x11 = -1 (0xFFFFFFFF), x9 = ~x10
                  // i.e., xor x9, x10, -1 = not x9, x10
```

x10	0000	0000	0000	0000	0000	1101	1100	0000
-----	------	------	------	------	------	------	------	------

x11	1111	1111	1111	1111	1111	1111	1111	1111
-----	------	------	------	------	------	------	------	------

x9	1111	1111	1111	1111	1111	0010	0011	1111
----	------	------	------	------	------	------	------	------

NOT Operations

- Useful to invert bits in a word

- ❖ Change 0 to 1, and 1 to 0

- RISC-V has **xori**

- ❖ **a XORI -1 == NOT a**

xori rd, rs, -1

-1 1111 1111 1111 1111 1111 1111 1111 1111

rs 0000 0000 0000 0000 0011 1100 0000 0000

XOR

rd 1111 1111 1111 1111 1100 0011 1111 1111

- Pseudo-instruction: **not**

not x10, x11

↓

xori x10, x11, -1

Basic Shifting

- Shift types
 - ❖ Logical (or unsigned) shift
 - ❖ Arithmetic (or signed) shift
- Shift directions
 - ❖ Left (multiply by powers of 2)
 - ❖ Right (divide by powers of 2)
 - Take floor value if the result is not an integer
 - Floor value of X (or $\lfloor X \rfloor$) is the largest integer less than or equal to X
 - $5 \gg 1 = \lfloor 5/2 \rfloor = 2$
 - $-3 \gg 1 = \lfloor -3/2 \rfloor = -2$

Logical Shift

■ Logical shift left

- ❖ MSB: shifted out
- ❖ LSB: shifted in with a 0
- ❖ $11001011 \ll 1 = 10010110$
- ❖ $11001011 \ll 3 = 01011000$

■ Logical shift right

- ❖ MSB: shifted in with a 0
- ❖ LSB: shifted out
- ❖ $11001011 \gg 1 = 01100101$
- ❖ $11001011 \gg 3 = 00011001$

- Logical shifts are useful to perform multiplication or division of unsigned integer by powers of two

Arithmetic Shift

■ Arithmetic shift left

- ❖ MSB: shifted out.
 - Be aware of overflow/underflow
- ❖ LSB: shifted in with a 0
- ❖ $1100 \ll 1 = 1000 \quad // -4 * 2 = -8$
- ❖ $1100 \ll 3 = 0000 \quad // -4 * 8 = 0 \text{ ??}$

■ Arithmetic shift right

- ❖ MSB: retain the sign bit
- ❖ LSB: shifted out
- ❖ $0100 \gg 1 = 0010 \quad // 4 / 2 = 2$
- ❖ $1100 \gg 1 = 1110 \quad // -4 / 2 = -2$

- Logical shifts are useful to perform multiplication or division of signed integer by powers of two
- The result of arithmetic shift left is the same as logical shift left

RISC-V Instructions – Shift operations (1)

■ **slli, srli**

❖ e.g., **slli rd, rs1, 4** ($rd \leftarrow rs1 \ll 4$)

■ **srai**

- ❖ Shift right arithmetic – extend the sign bit
- ❖ Useful for dividing 2's complement numbers

RISC-V Instructions – Shift operations (2)

- **sll, srl, sra**

- ❖ e.g., **sll rd, rs1, rs2** ($rd \leftarrow rs1 \ll rs2$)

- Shift by the amount held in the lower 5 bits of register **rs2**

```
li x10, 200
```

```
li x11, 100    //  $100_{10} = 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0110\ 0100_2$ 
```

```
sll x12, x10, x11    //  $x12 = 200 \ll 00100_2 = 200 * 2^4 = 3200$ 
```

RISC-V Instructions – Shift operations (3)

- $-123 / 2^4 = -123 / 16 = -7.6875 = -7 ? -8?$

❖ Shift right arithmetic: “round down”

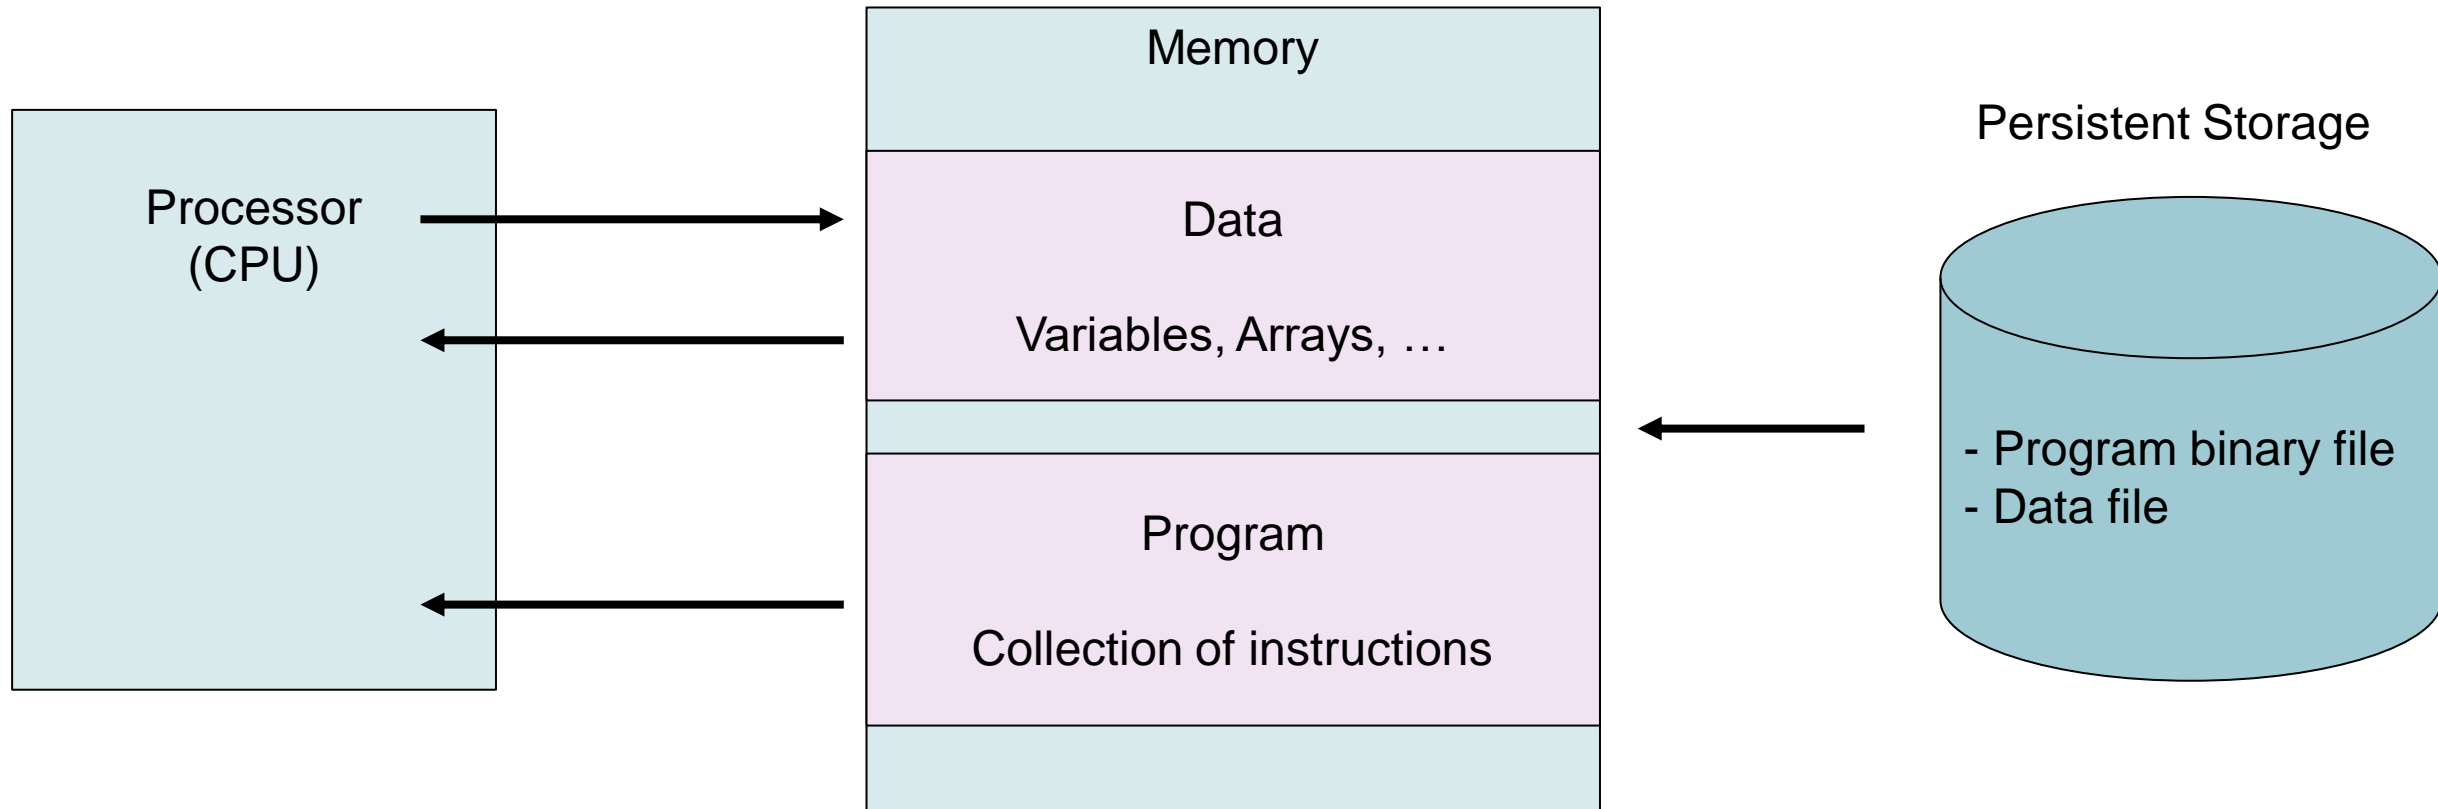
`li x10, -123` // $-123_{10} = 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1000\ 0101_2$

`li x11, 4`

`sra x12, x10, x11` // $x12 = 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1000_2$
 $= -8_{10}$

MACHINE CODE

Stored Program Concept




- Instructions and data are represented in binary
- Instructions and data are stored in memory
- CPU fetches instructions and data to execute

Representing Instructions (1)

- Instructions are encoded in binary
 - ❖ Called machine code
- RISC-V instructions
 - ❖ Encoded as 32-bit (4 bytes) instruction words
 - ❖ CISC processors have variable instruction length
 - x86: 1byte – 15 bytes


RISC-V R-format Instructions



■ Instruction fields

- ❖ **opcode**: operation code
- ❖ **rd**: destination register number
- ❖ **funct3**: 3-bit function code (additional opcode)
- ❖ **rs1**: the first source register number
- ❖ **rs2**: the second source register number
- ❖ **funct7**: 7-bit function code (additional opcode)

RISC-V R-format Instructions




add x9, x20, x21

funct7	x21	x20	funct3	x9	Arithmetic
0	21	20	0	9	51
0000000	10101	10100	000	01001	0110011

0000000 10101 10100 000 01001 0110011₂
0000 0001 0101 1010 0000 0100 1011 0011₂
= 0x015A04B3

RISC-V R-format Instructions




sub x9, x20, x21

funct7	x21	x20	funct3	x9	Arithmetic
32	21	20	0	9	51
0100000	10101	10100	000	01001	0110011


0100000 10101 10100 000 01001 0110011₂
0100 0001 0101 1010 0000 0100 1011 0011₂
= 0x415A04B3

RISC-V I-format Instructions



- Immediate arithmetic and load instructions
 - ❖ **rs1**: source or base address register number
 - ❖ **immediate**: constant operand, or offset added to base address
 - 2's-complement
 - Will be sign extended to 32-bit

RISC-V I-format Instructions




addi x9, x20, 123

immediate	x20	funct3	x9	Imm. Arithmetic
123	20	0	9	19
0000 0111 1011	10100	000	01001	0010011

$000001111011\ 10100\ 000\ 01001\ 0010011_2$
 $0000\ 0111\ 1011\ 1010\ 0000\ 0100\ 1001\ 0011_2$
 = 0x7ba0493

RISC-V I-format Instructions



andi x9, x20, -123

immediate	x20	funct3	x9	Imm. Arithmetic
-123	20	7	9	19
1111 1000 0101	10100	111	01001	0010011

111110000101 10100 111 01001 0010011₂
1111 1000 0101 1010 0111 0100 1001 0011₂
= 0xF85A7493

Sign Extension

- Representing a number using more bits
 - ❖ Preserve the numeric value
- Replicate the sign bit to the left
- Examples: 8-bit to 16-bit
 - ❖ +2: 0000 0010 => 0000 0000 0000 0010
 - ❖ -2: 1111 1110 => 1111 1111 1111 1110
- In RISC-V instruction set,
 - ❖ lb: sign-extend loaded byte
 - ❖ lbu: zero-extend loaded byte

Sign Extension For Logic Operations

li x20, 0x12345 // x20 = 0000 0000 0000 0001 0010 0011 0100 0101₂

andi x9, x20, -123 // -123 = 1111 1000 0101₂

// → 1111 1111 1111 1111 1111 1111 1000 0101₂


x20	0000	0000	0000	0001	0010	0011	0100	0101
-123	1111	1111	1111	1111	1111	1111	1000	0101
x9	0000	0000	0000	0001	0010	0011	0000	0101

0x12305

Shift Operations

■ slli, srli, srai

- ❖ shamt: shift amount in number of bit positions



■ sll, srl, sra

- ❖ $rd \leftarrow$ value in rs1 shifted by shift amount in rs2[4:0]

