

Introduction to Computer Architecture

Chapter 2

Instructions: Language of the Computer - 3

Hyungmin Cho

Department of Computer Science and Engineering
Sungkyunkwan University

Procedure Calling

Caller

```
int main() {  
    ...  
    ...  
    arg2 = val1 * 2;  
    result = compute(arg1, arg2);  
    ...  
    ...  
}
```

Callee

```
int compute(int p1, int p2) {  
    int temp1;  
    int temp2;  
    ...  
    temp1 = p1 * p2;  
    temp2 = temp1 + 10;  
    ...  
    return temp2;  
}
```

1. How to pass parameters (arg1 and arg2) to the **callee** (function)
2. How to reserve memory space for the callee's **local variables**
3. How to transfer the result (return value) back to the **caller**
4. How to return to the **caller**

Procedure? Function?

- Are they the same? ... Not really.
- Procedure is a group of instructions that are invoked to perform a designated task.
- There is no concept of “Function” in assembly language
- We use a special rule to support “C”-style function in assembly
→ calling convention

RISC-V Register Usage

Register Number	Name	Usage
x0	zero	Constant 0 (hardwired)
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7, x28-x31	t0 - t6	Temporaries
x8	s0 / fp	Frame pointer
x9, x18-x27	s1 - s11	Saved registers
x10-x11	a0 - a1	Function arguments / results
x12-x17	a2 - a7	Function arguments

Argument Passing

- Set register $x10 - x17$ with the arguments in order **before calling**
 $(a0) \quad (a7)$
the function.
- If the function takes more than 8 arguments, the rest of the arguments are passed to the callee using memory.

Return Value

- Use register x10 – x11
 - ❖ x10 [31:0]
 - ❖ x11 [63:32]
- If the return value is larger than 64-bit, the return value is passed through memory.

Procedure Call Instructions

- Procedure call: jump-and-link instruction

jal x1, ProcedureLabel

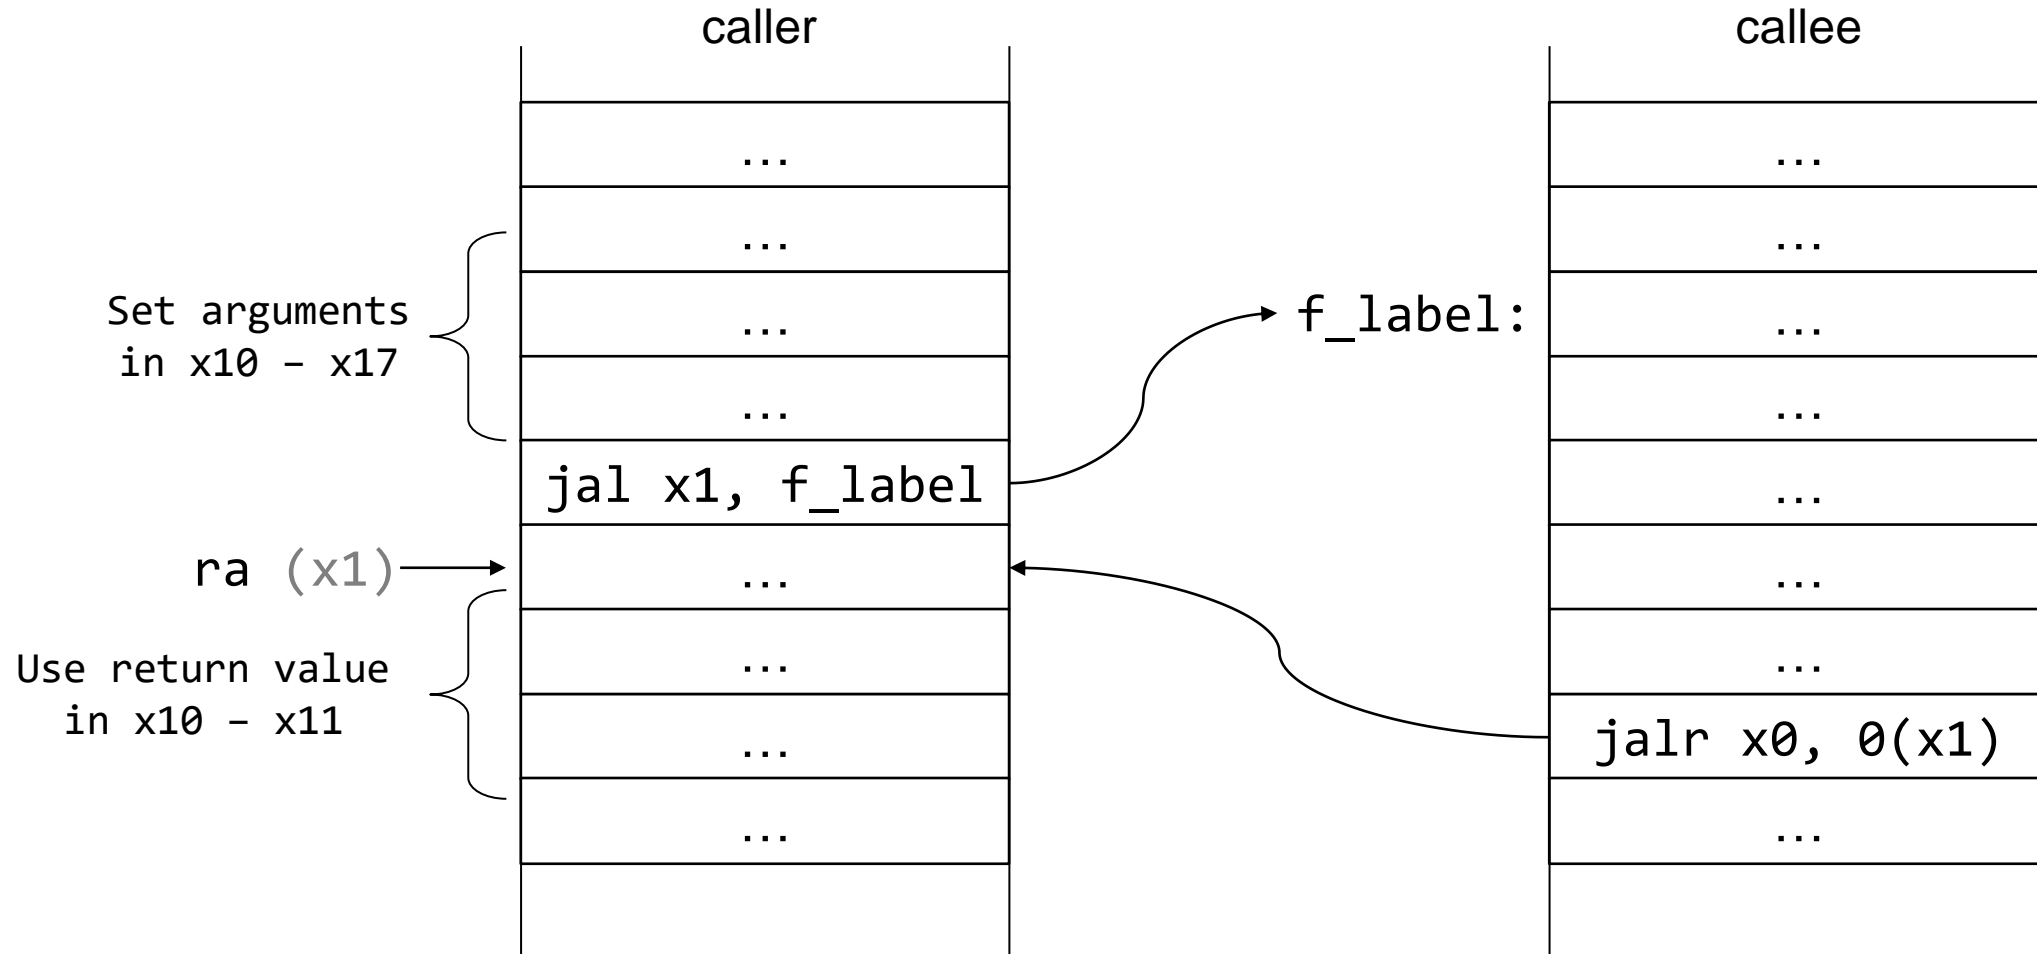
- ❖ Jumps to the target address
- ❖ Save the address of the next instruction in **x1**
- ❖ **x1**'s register name is **ra** (return address)

- Procedure return: jump-and-link register instruction

jalr x0, Offset(x2)

- ❖ Jump to offset + address in x2
- ❖ Address of the next instruction → x0
- ❖ Can also be used for computed jumps
 - e.g., for case/switch statements

RISC-V Procedure Call



Memory Space for Procedure (Function)

- A procedure needs its own memory space for the the local variables.

```
int func1(...) {  
    int var1;  
    int arr2[2];  
    ...  
}
```

- ❖ If all local variables can fit in the registers, a procedure may not need to use memory space (compiler optimization)
- ❖ Even all local variables can fit in the registers, the values may need to be saved in the memory to call another function (procedure).

```
int func1(...) {  
    int var1;  
    var1 = 123  
    ...  
    func2(...)  
    var2 = var1*2  
    ...  
}
```

x5 ←

```
int func2(...) {  
    int var10;  
    ...  
    var10 = 456;  
    ...  
}
```

→ x5

Memory Space for Procedure (Function)

- How to reserve space for each procedure?
 - ❖ Fixed location for each procedure **is not an ideal solution**
 - Can't know which procedures will be called in a program in advance
 - Recursion!


Fixed address?

```
int func1(...) {  
    int var1;  
    int arr2[2];  
    ...  
}
```

→ 0x1000
→ 0x1004

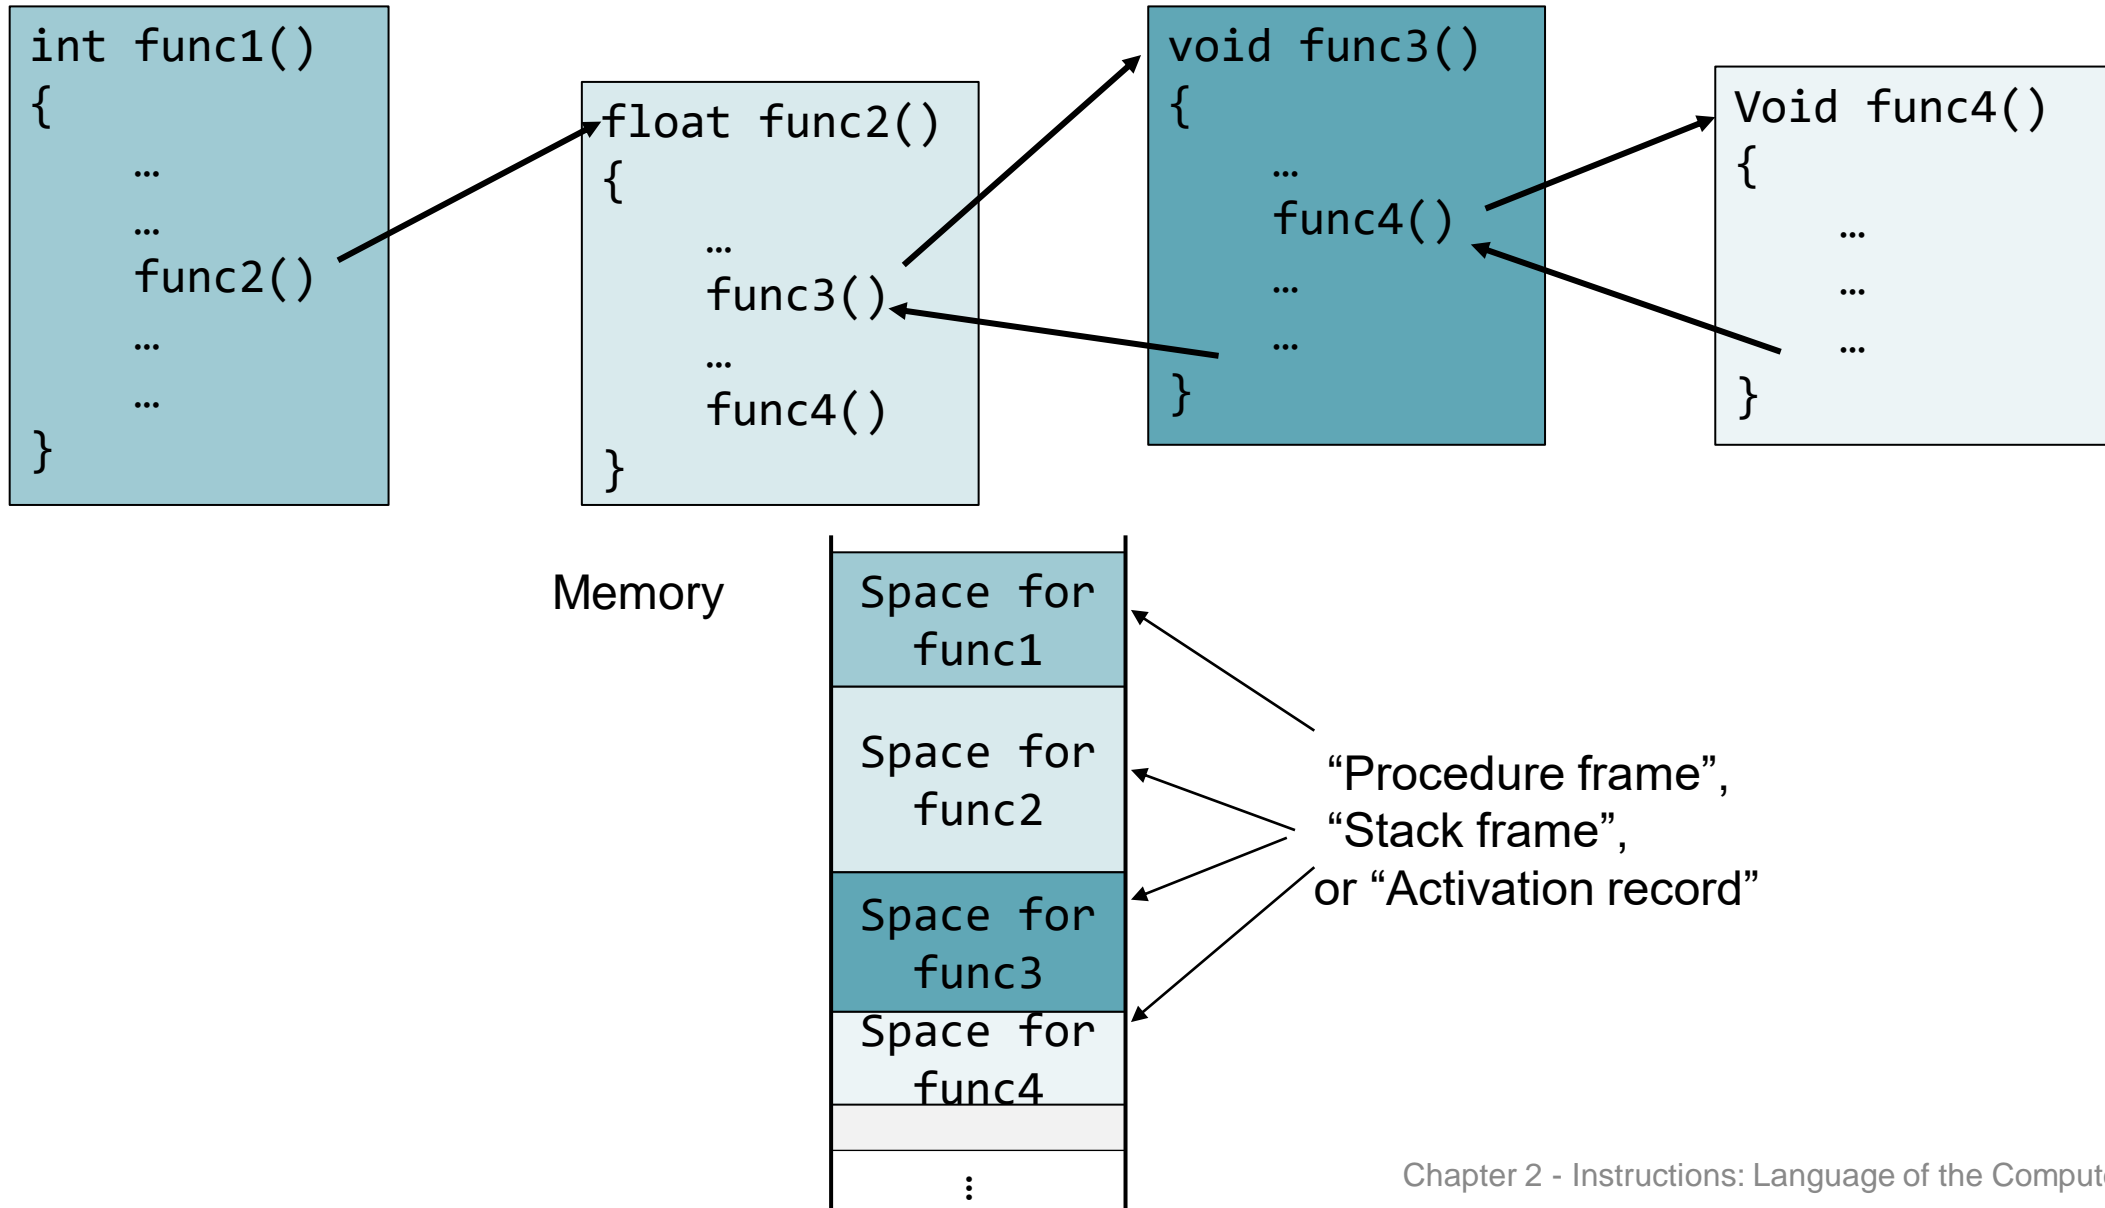
```
int func2(...) {  
    int arr20[100];  
    int var10;  
    ...  
}
```

→ 0x2000
→ 0x2190

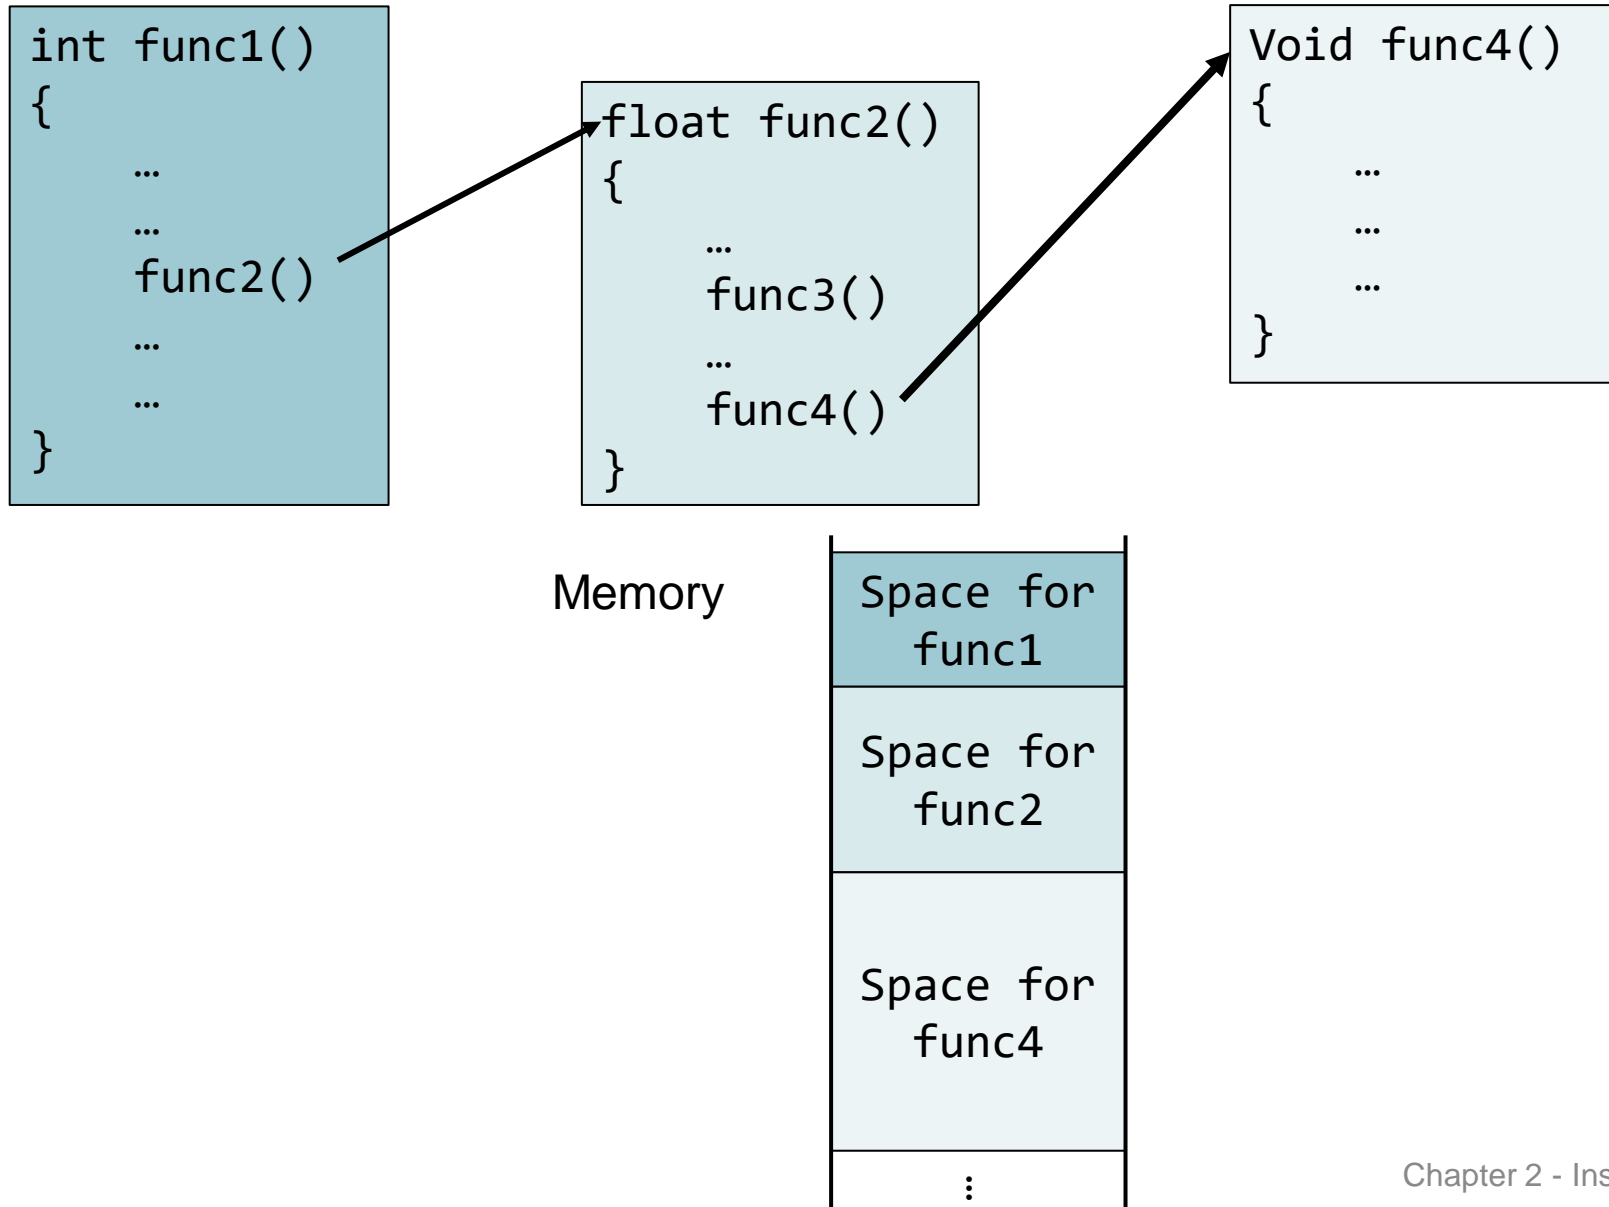


- Solution?
 - ❖ “Allocate” the space when the function is called / “Deallocate” when returns
 - ❖ Use relative addressing within the allocated space.

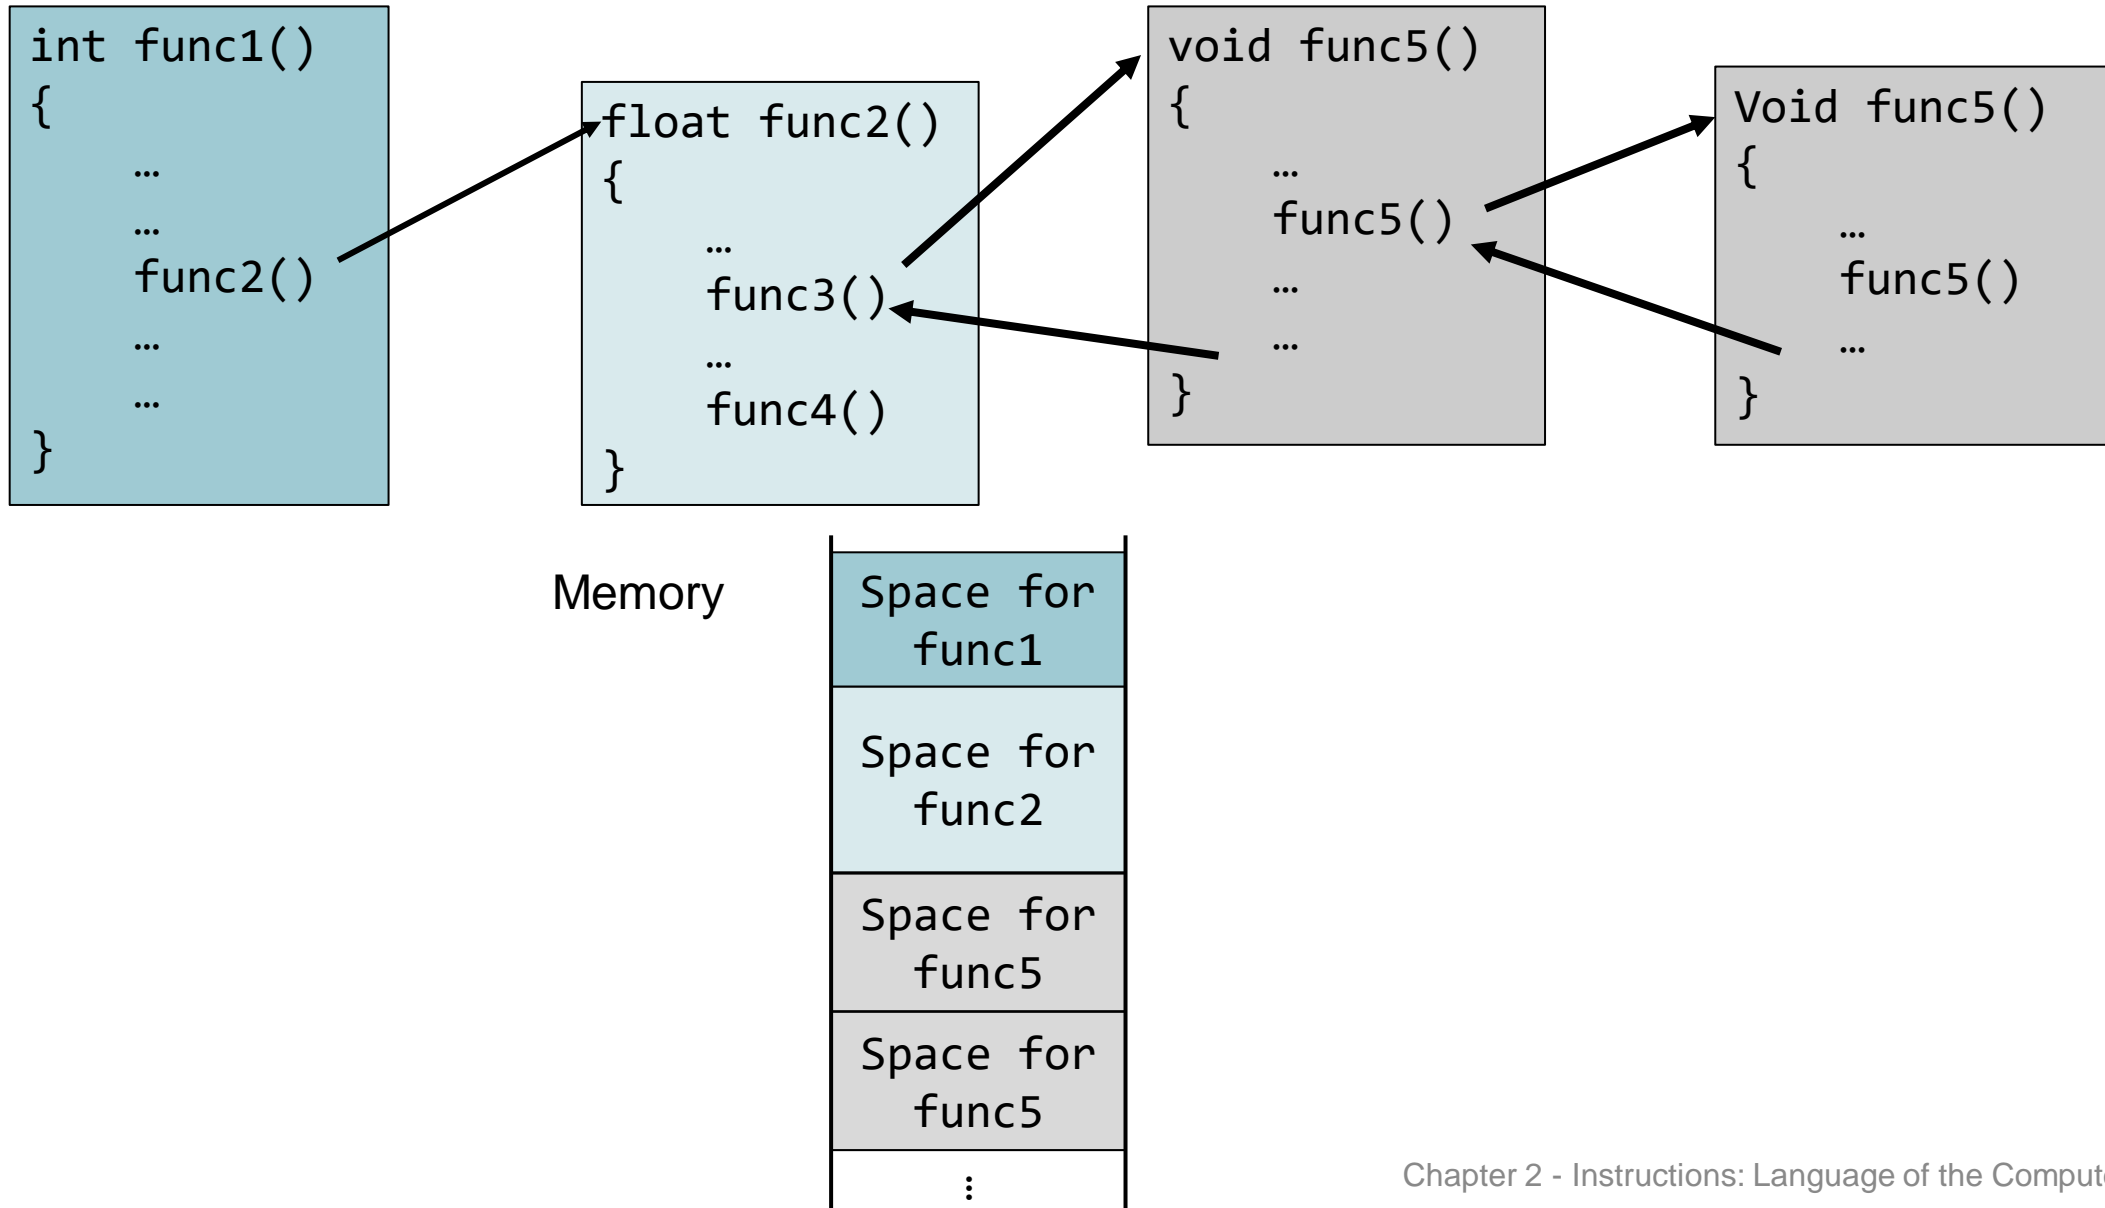
Stack



Stack




Stack



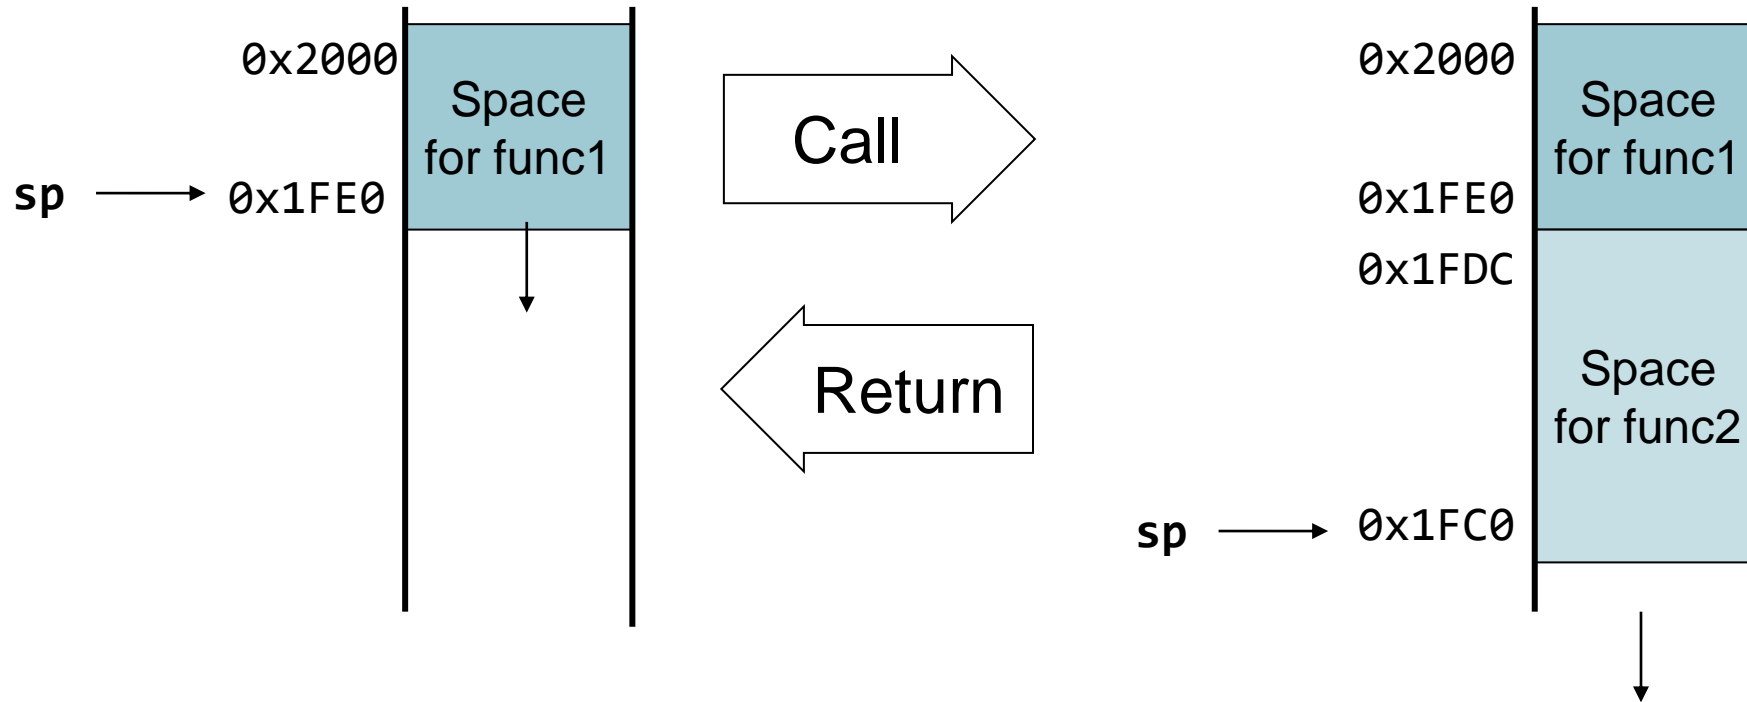
Managing & Accessing Stack (1)

- A stack frame can be placed at any location in memory
 - ❖ Relative addressing is required
- Stack pointer register x2 (sp) points to the lowest address in the current stack frame



Managing & Accessing Stack (2)

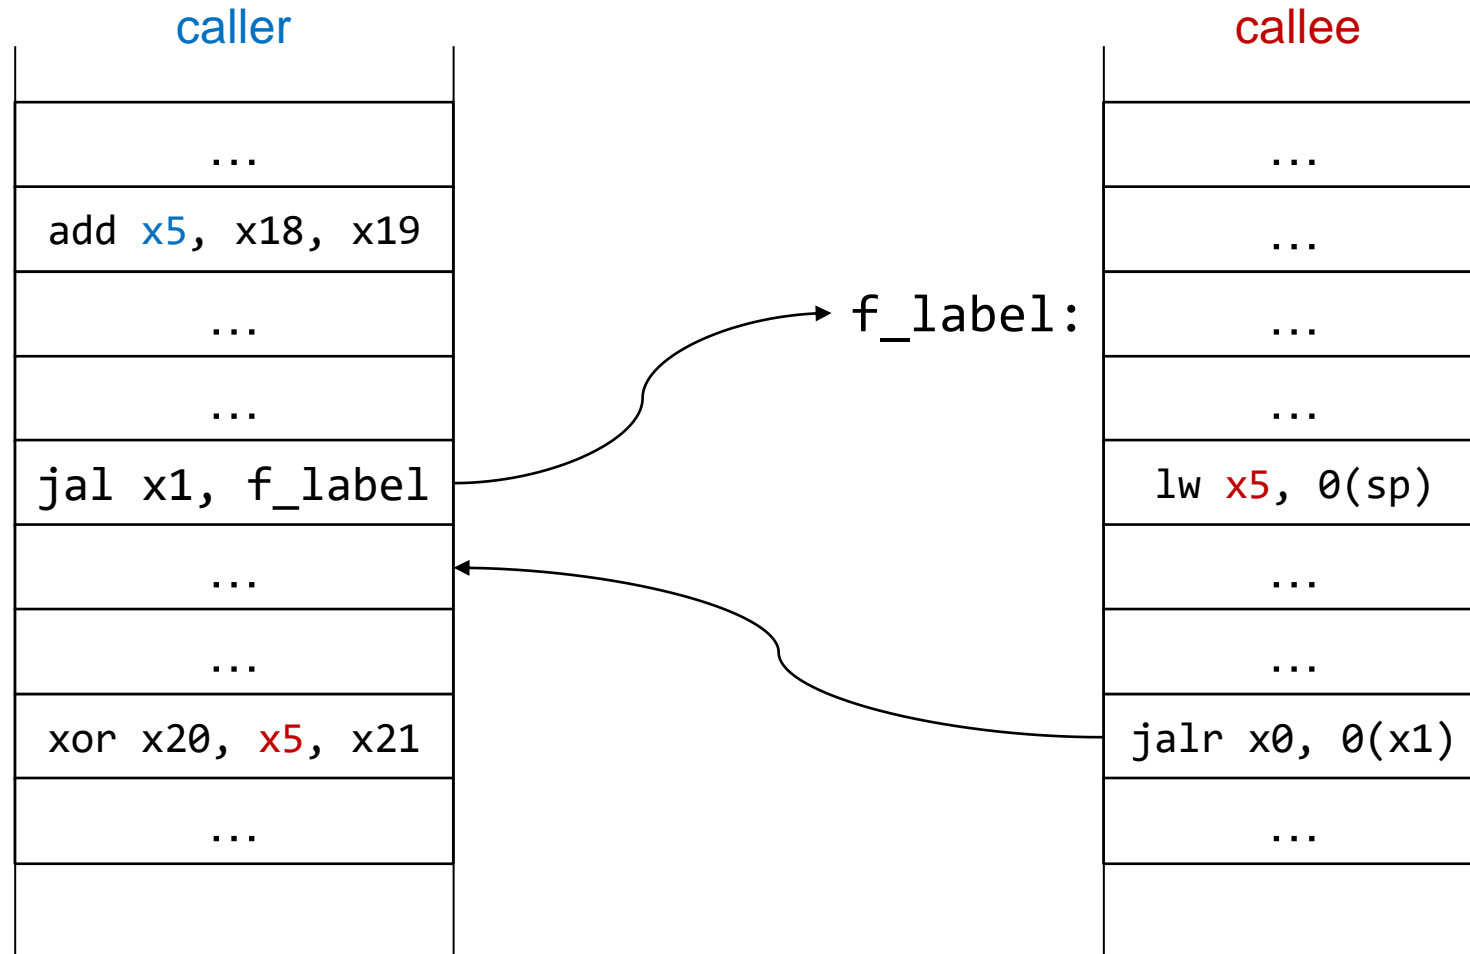
- Stack grows from higher address to lower address



Managing & Accessing Stack (3)

- Each procedure needs to know how much stack space it needs
 - ❖ High-level language: the compiler automatically computes the required space
 - ❖ Assembly language: the programmer need to determine the required size
- Suppose a function requires 20 bytes of space (e.g., 5 x 4-byte values)...
 - ❖ When the function starts, **decrement** `sp` by 20
 - `addi sp, sp, -20`
 - ❖ To access any stack value, use `sp` as the base register
 - `sw x18, 4(sp)`
 - ❖ When the function returns, **increment** `sp` by 20
 - `addi sp, sp, 20`

Saving Registers



Register Saving Convention

Register Number	Name	Usage
x0	zero	Constant 0 (hardwired)
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-x7, x28-x31	t0 - t6	Temporaries
x8	s0 / fp	Frame pointer
x9, x18-x27	s1 - s11	Saved registers
x10-x11	a0 - a1	Function arguments / results
x12-x17	a2 - a7	Function arguments

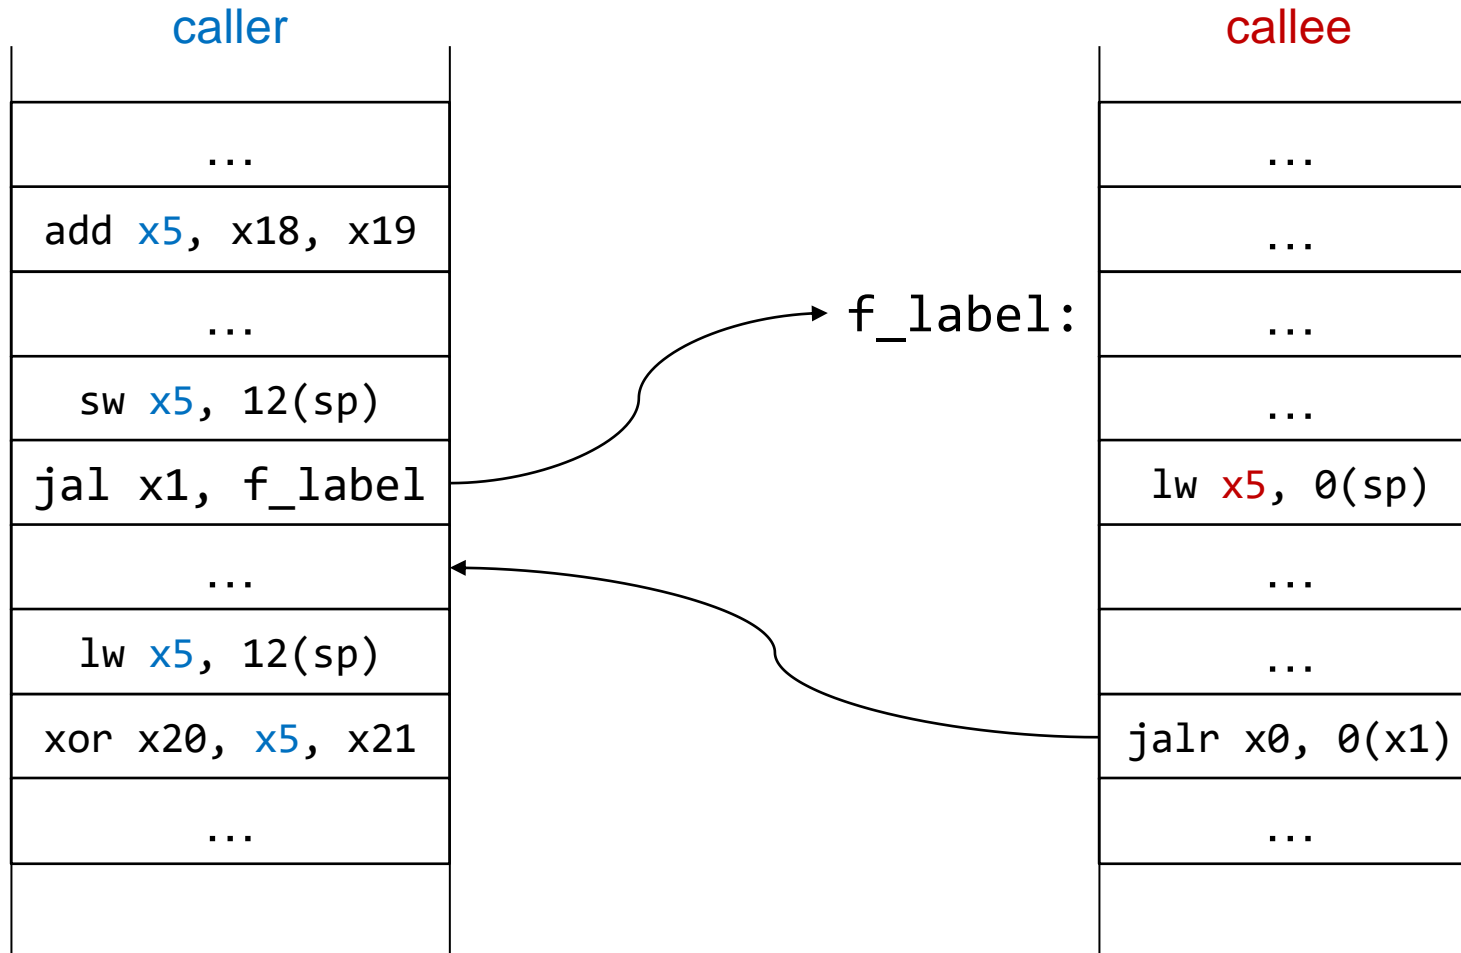
Caller-save registers

Callee-save registers

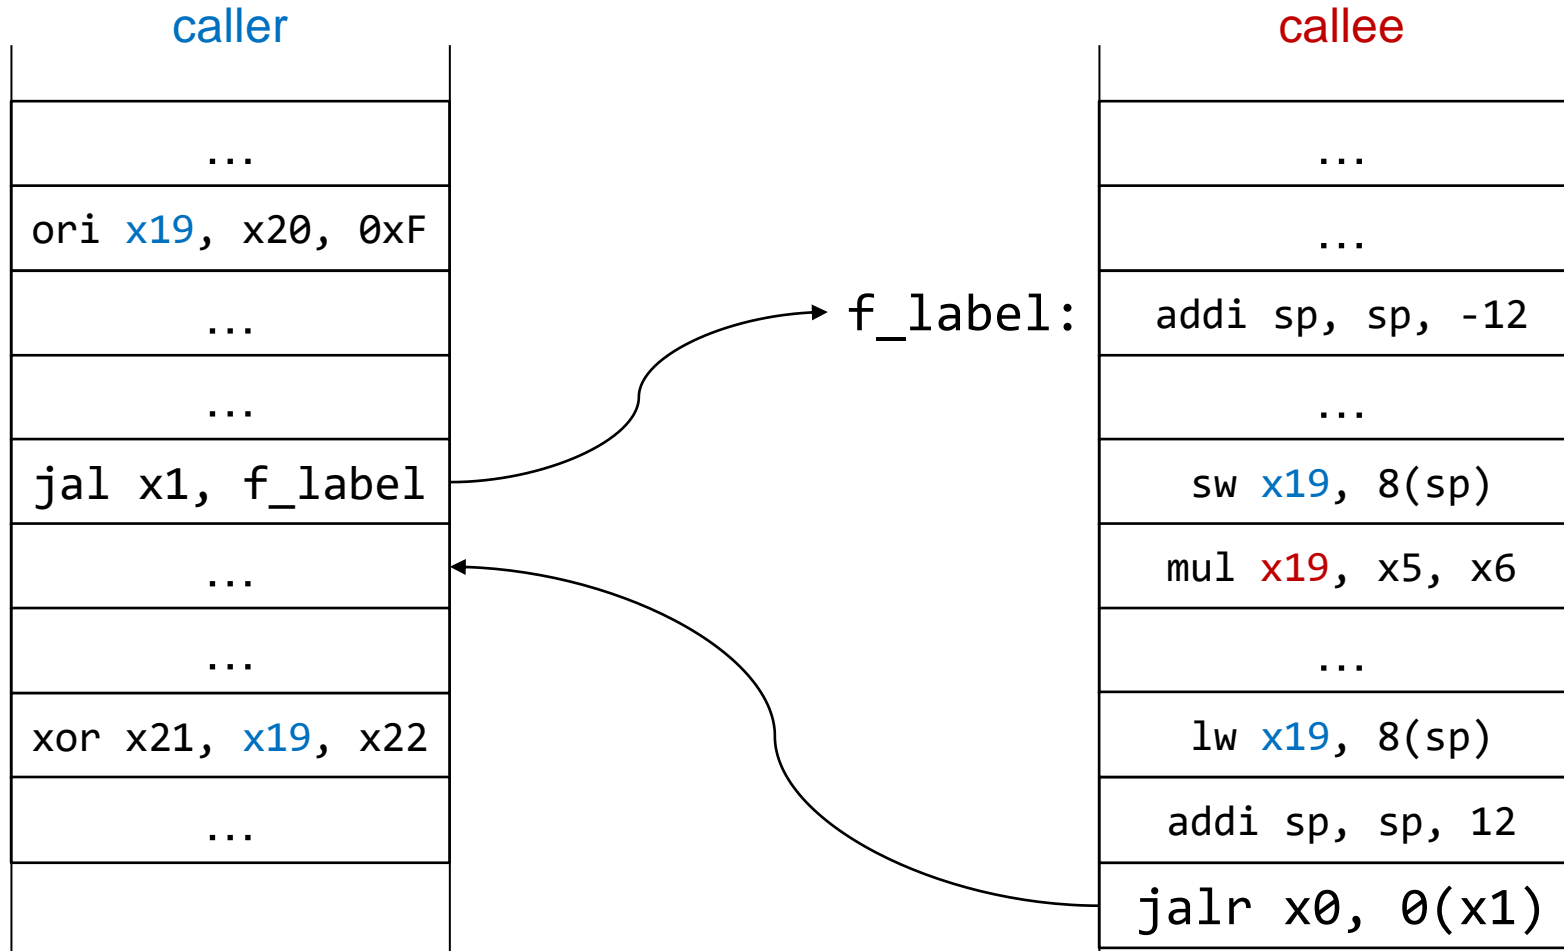
If the caller wants to **use value in these registers** *after* calling a function,
→ caller need to store these registers in its stack frame before calling the function

If the callee wants to **use these registers**,
→ callee need to store these registers in its stack frame and restore those before the return

Caller-save Registers



Callee-save Registers



Avoiding Unnecessary Register Save & Restore

- If a value is used for a short period, use the “Temporary” registers
(caller-save)
instead of the “Saved” registers
(callee-save)
- Use “Saved” registers if a value is created before a function call and used after the function call.
- Use “Argument” and “Results” registers as temporary registers if possible

Leaf Procedure Example

- C code:

```
int leaf_example (int g, h, i, j)
{
    int f;
    f = (g + h) - (i + j);
    return f;
}
```

- ❖ Arguments g, ..., j in x10, x11 ..., x13 (a0, a1, ..., a3)
- ❖ f in x20 (s4)
 - hence, need to save the caller's x20 on stack
- ❖ Temporaries x5 (t0), x6 (t1)
- ❖ Return value in x10 (a0)

Leaf Procedure Example

```
int leaf_example (int g, h, i, j)
{
    int f;
    f = (g + h) - (i + j);
    return f;
}
```

leaf_example:

```
addi sp, sp, -4
sw    x20, 0(sp)
add   x5, x10, x11
add   x6, x12, x13
sub   x20, x5, x6
addi  x10, x20, 0
lw    x20, 0(sp)
addi  sp, sp, 4
jalr  x0, 0(x1)
```

Allocate stack space

Save x20 on stack

} Procedure body

Set return value

Restore x20

De-allocate stack space

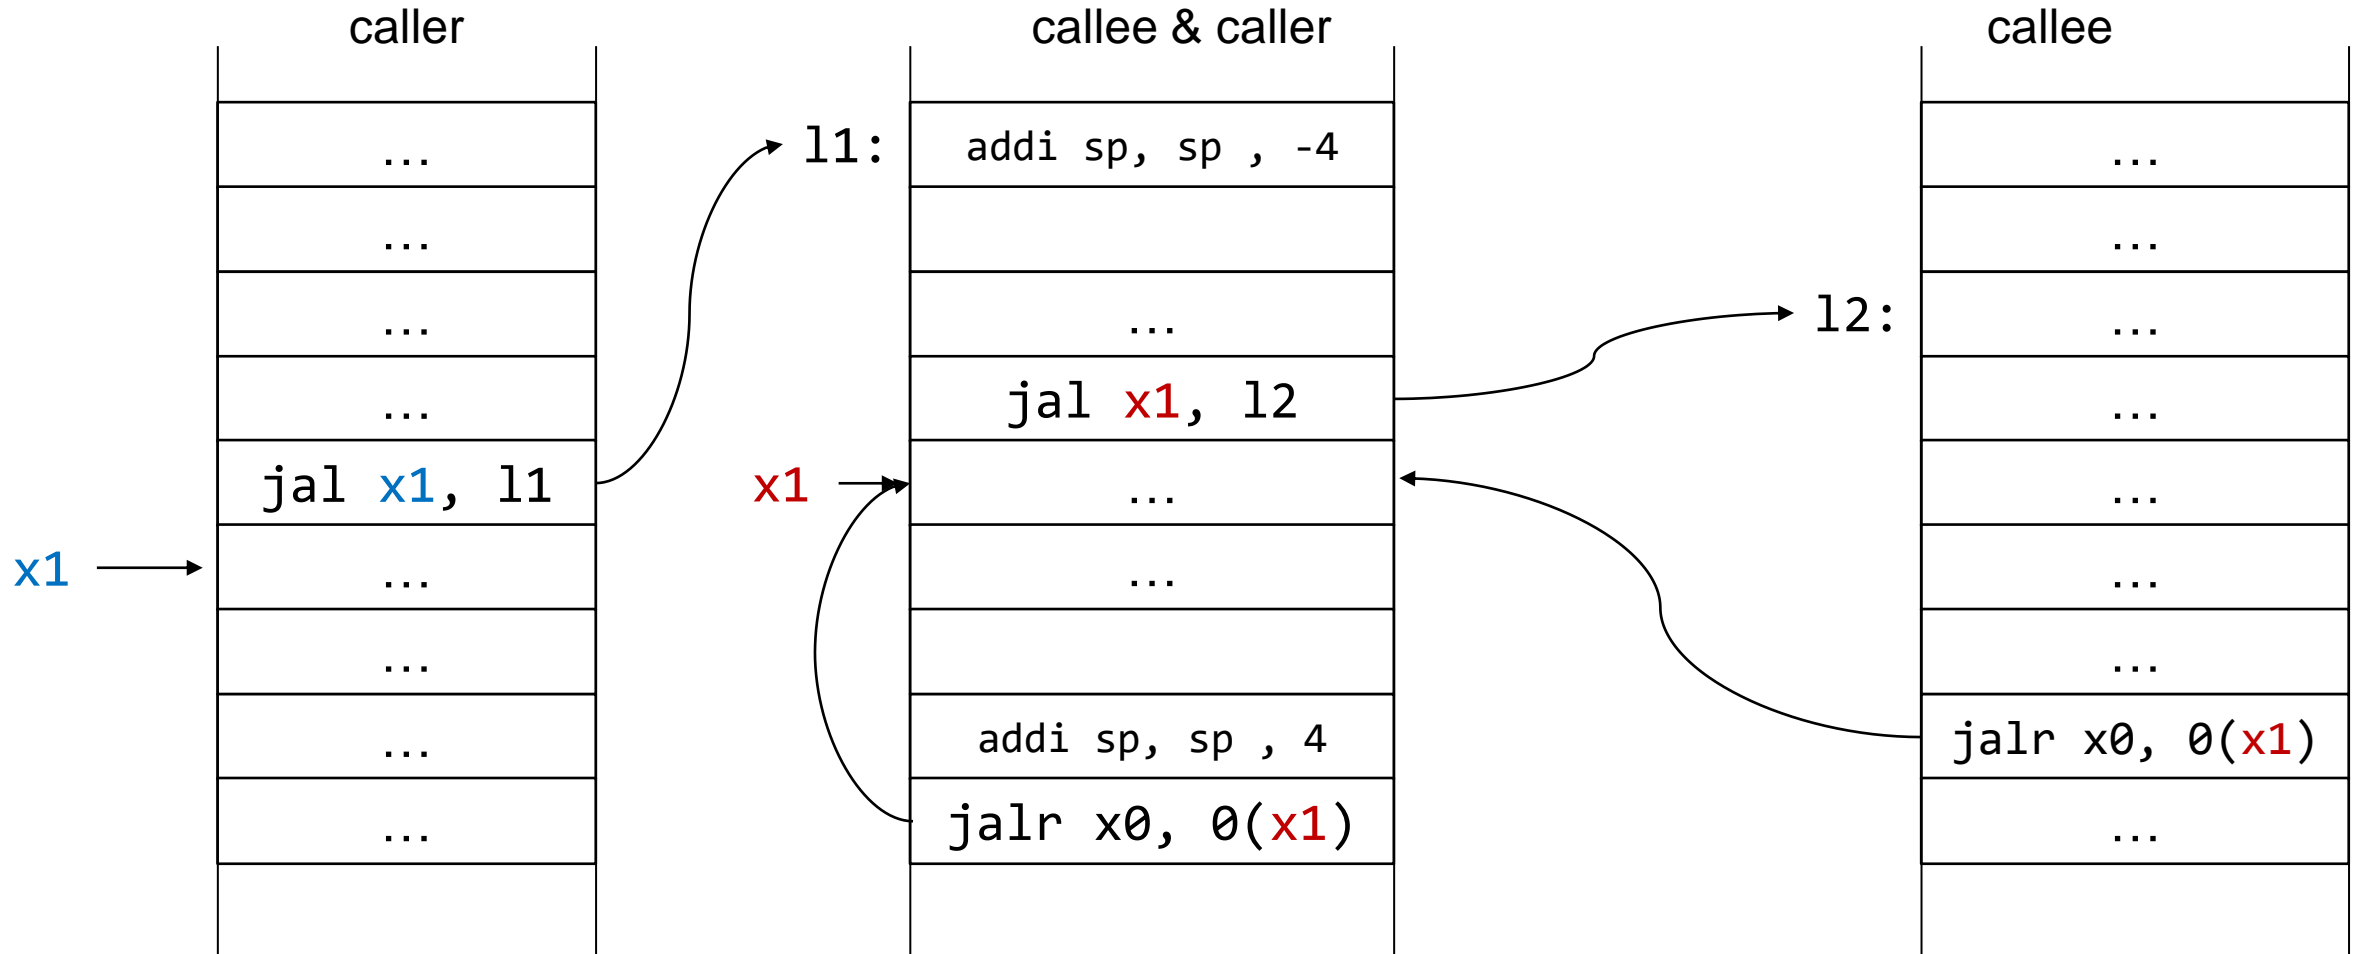
Return

Leaf Procedure Example (2)

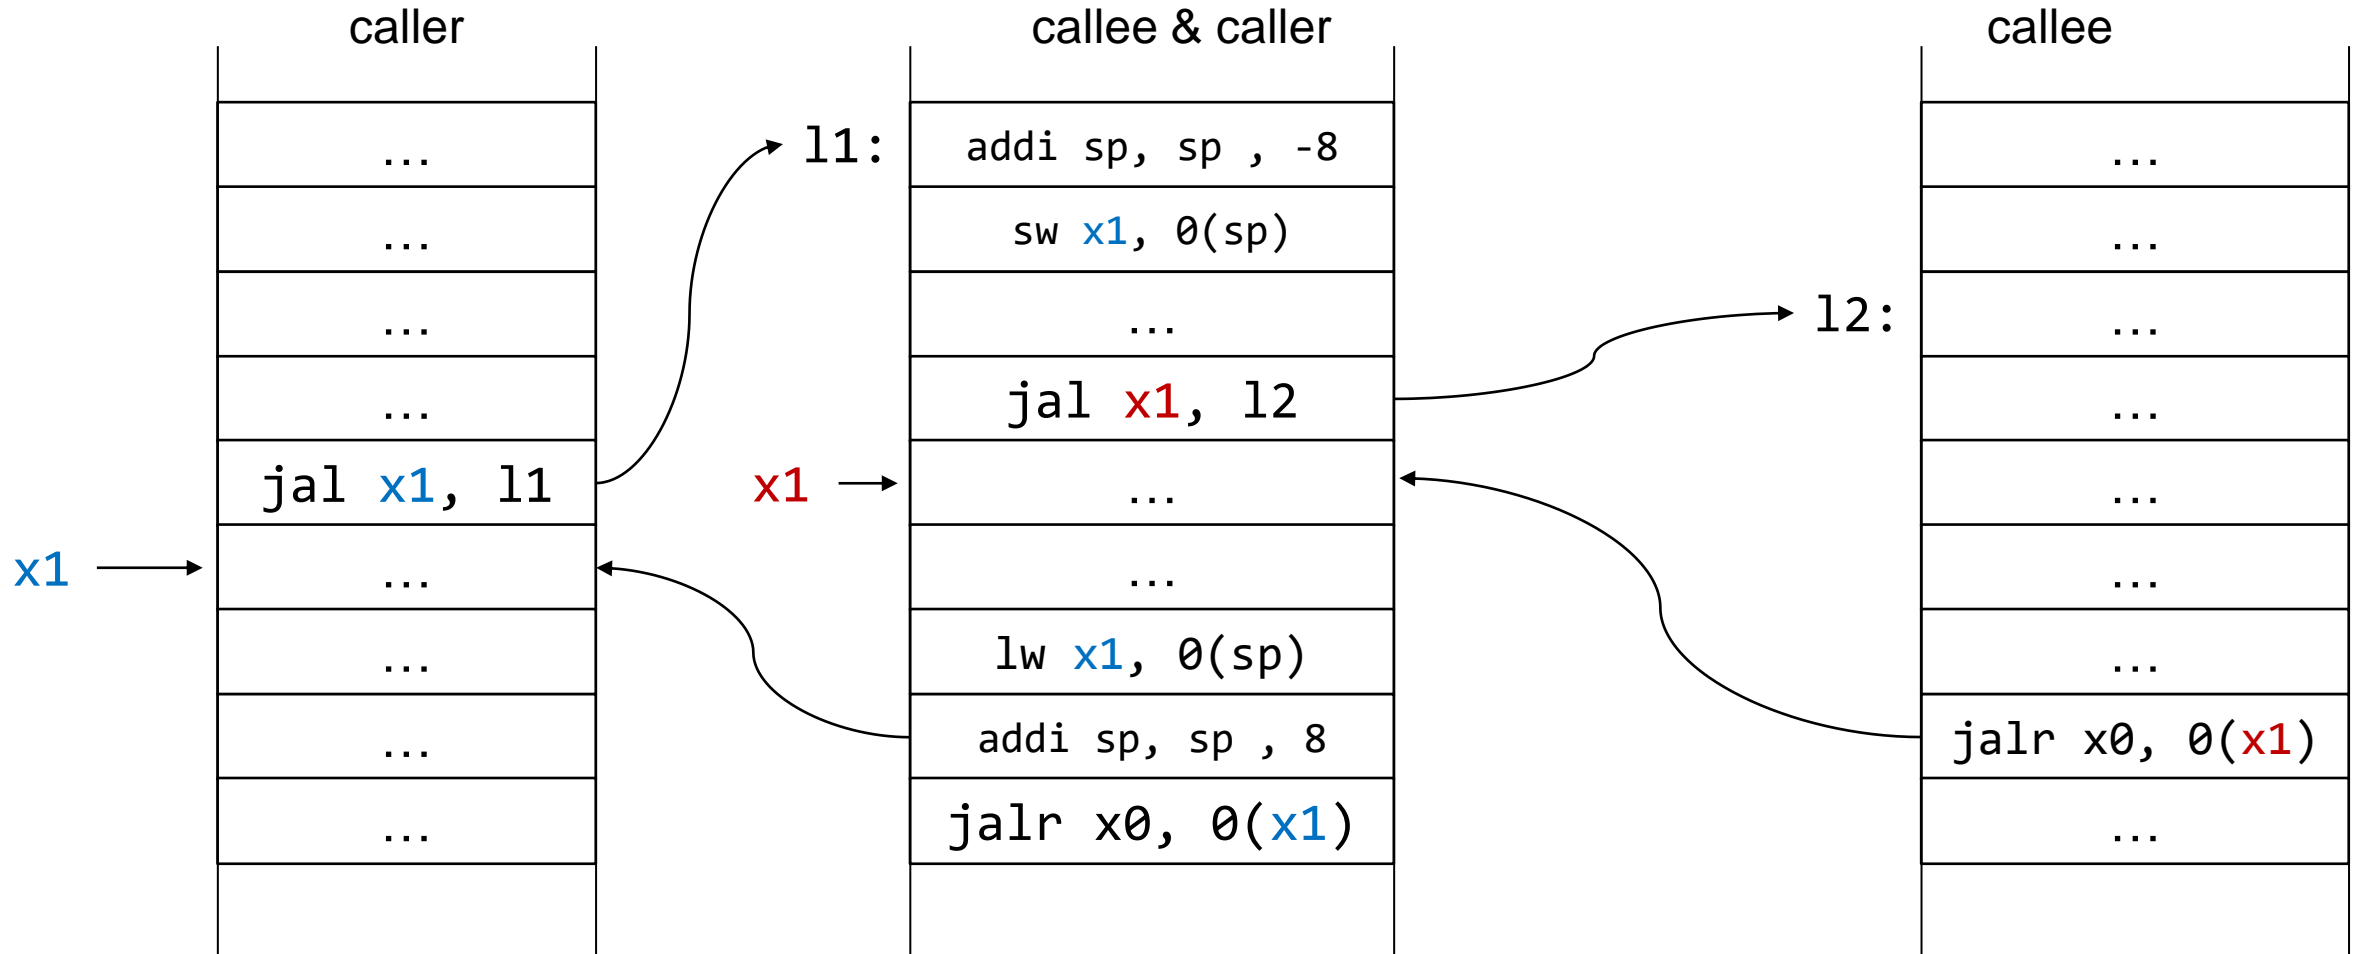
```
int leaf_example (int g, h, i, j)
{
    int f;
    f = (g + h) - (i + j);
    return f;
}
```

```
leaf_example:
    add    x5, x10, x11
    add    x6, x12, x13
    sub    x10, x5, x6
    jalr   x0, 0(x1)
```


MIPS Procedure Calls – Non-Leaf



MIPS Procedure Calls – Non-Leaf



Non-Leaf Procedure Example

- C code:

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

- ❖ Argument n in x10
- ❖ Result in x10
- ❖ Need to store return address (x1) in stack
- ❖ Need to store argument (x10) in stack
 - n is used again after calling fact(n-1)

Non-Leaf Procedure Example

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

fact:

```
addi sp, sp, -8      # allocate stack for two 4-byte items
sw    x1, 4(sp)      # save return address
sw    x10, 0(sp)     # save argument (n)
```

```
addi x5, x0, 1
bge  x10, x5, L1      # test if n < 1
```

```
addi x10, x0, 1      # if n < 1, result is 1
addi sp, sp, 8        # de-allocate stack space
jalr x0, 0(x1)        # return
```

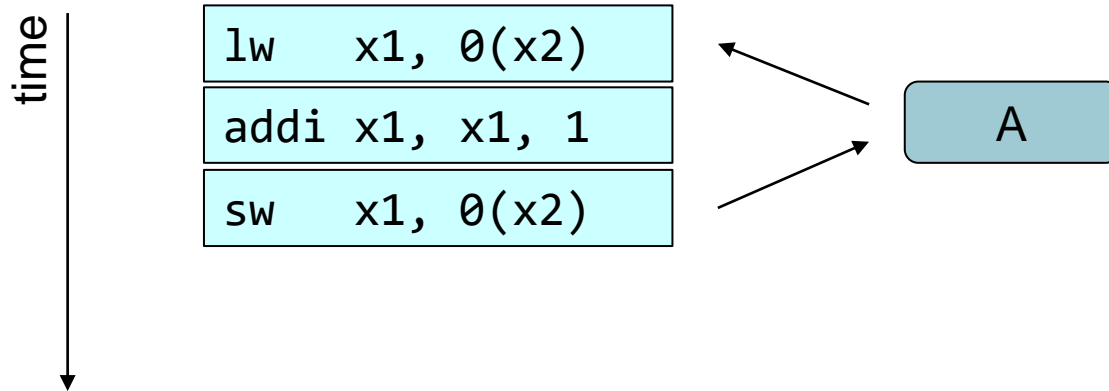
```
L1: addi x10, x10, -1 # if n >= 1 set n-1 as argument
jal  x1, fact        # recursive call
```

```
lw    x5, 0(sp)      # restore original n
mul   x10, x5, x10    # multiply to get result
```

```
lw    x1, 4(sp)      # restore return address
addi sp, sp, 8        # de-allocate stack space
jalr x0, 0(x1)        # return
```

Synchronization

- $A = A + 1;$




Synchronization

■ $A = A + 1;$

time
↓

lw	x1, 0(x2)
addi	x1, x1, 1
sw	x1, 0(x2)



■ $A = A + 2;$

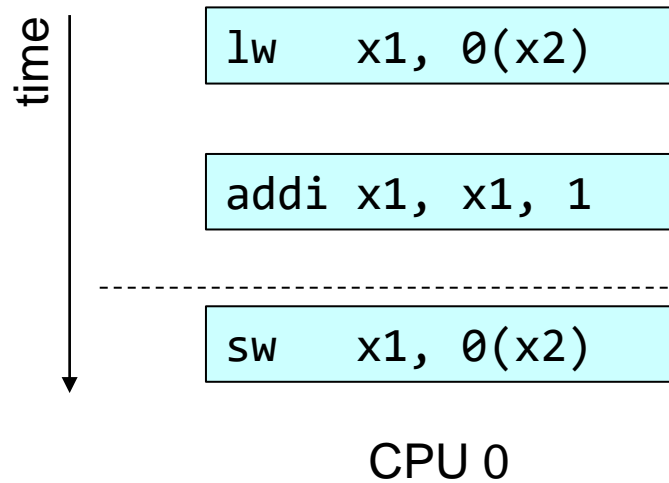
lw	x1, 0(x2)
addi	x1, x1, 2
sw	x1, 0(x2)

CPU 1


$A = A + 3$

Synchronization: Race Condition

■ $A = A + 1;$



■ $A = A + 2;$




$A = A + 1$??

Data "race"

Synchronization: Atomic Operation (1)

- Special memory instructions (RISC-V “A” extension)



- ❖ lr.w : load word **reserved**
 - Load from address in x2 to x1 (no offset)
- ❖ sc.w : store word **conditional**
 - Store from x1 to address in x2 (no offset)
- sc.w writes to the memory only if the target location wasn't modified after lr.w
 - ❖ → x1 is set to 0
- If the memory location was changed by someone else, store does not happen
 - ❖ → x1 is set to a non-zero value

Synchronization: Atomic Operation (2)

■ $A = A + 1;$

```
lr.w  x1, (x2)
```

```
addi x1, x1, 1
```

```
sc.w  x1, (x2)
```

CPU 0

- Write does not happen
- x1 becomes non-zero
- Your program may choose to retry the task if x1 is non-zero.

■ $A = A + 2;$

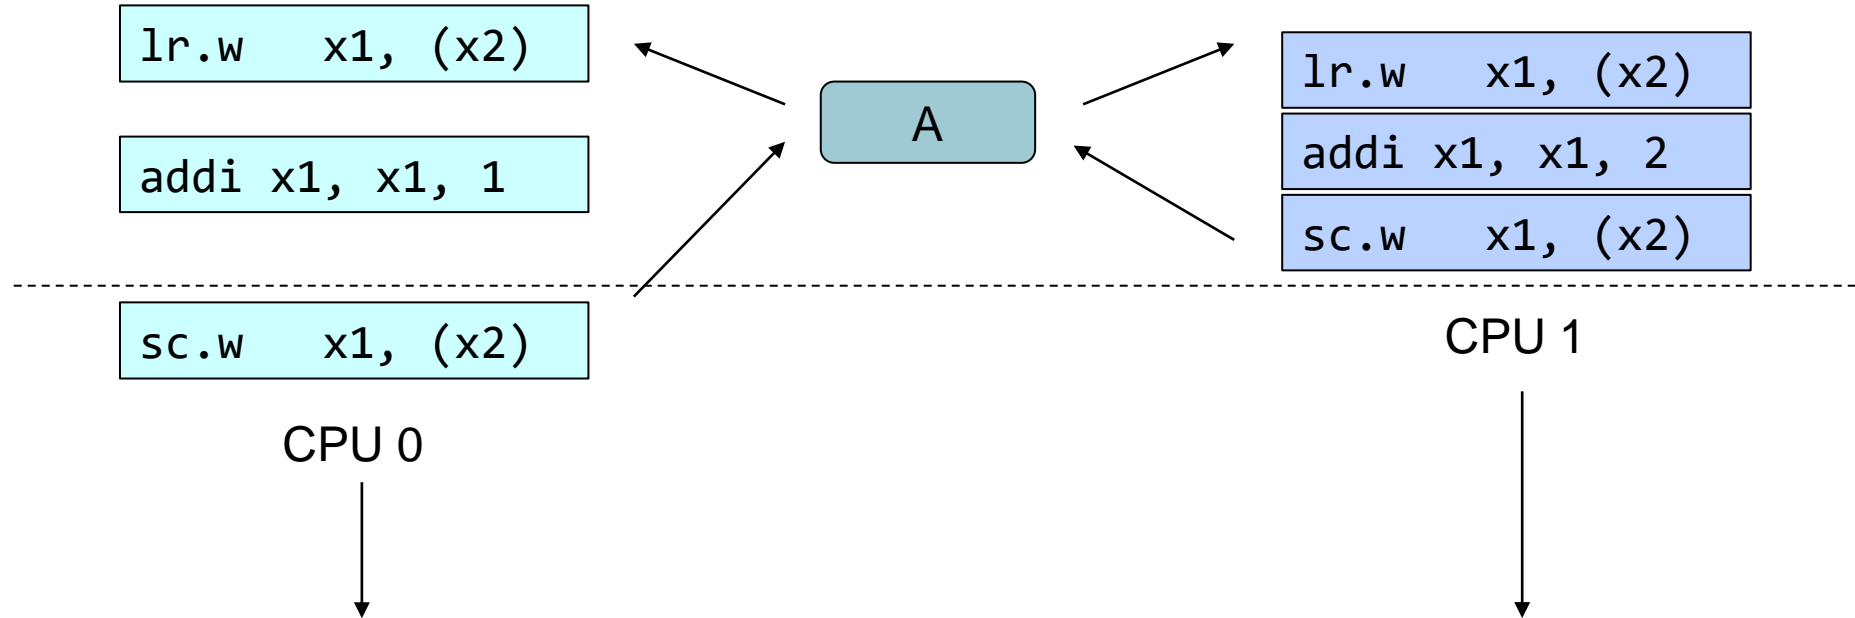
```
lr.w  x1, (x2)
```

```
addi x1, x1, 2
```

```
sc.w  x1, (x2)
```

CPU 1

- Write happens
- x1 becomes 0
- Your program can proceed to the next task if x1 is 0.



Synchronization

- In high-level languages, you can use these basic synchronization methods through library functions
- In Linux, POSIX pthread library
 - ❖ `pthread_mutex_lock()`
 - ❖ `pthread_mutex_unlock()`