

# **Introduction to Computer Architecture**

## **Chapter 5**

### **Large and Fast: Exploiting Memory Hierarchy - 2**

**Hyungmin Cho**

Department of Computer Science and Engineering  
Sungkyunkwan University

# Improving Cache Performance

---

## ■ Associativity

- ❖ Drawback of a direct-mapped cache:

- If frequently used addresses are mapped to the same cache index, many cache misses will happen!

→ Solution: Multiple slots per index!

## ■ Multi-level Cache

- ❖ Small cache has high miss rate

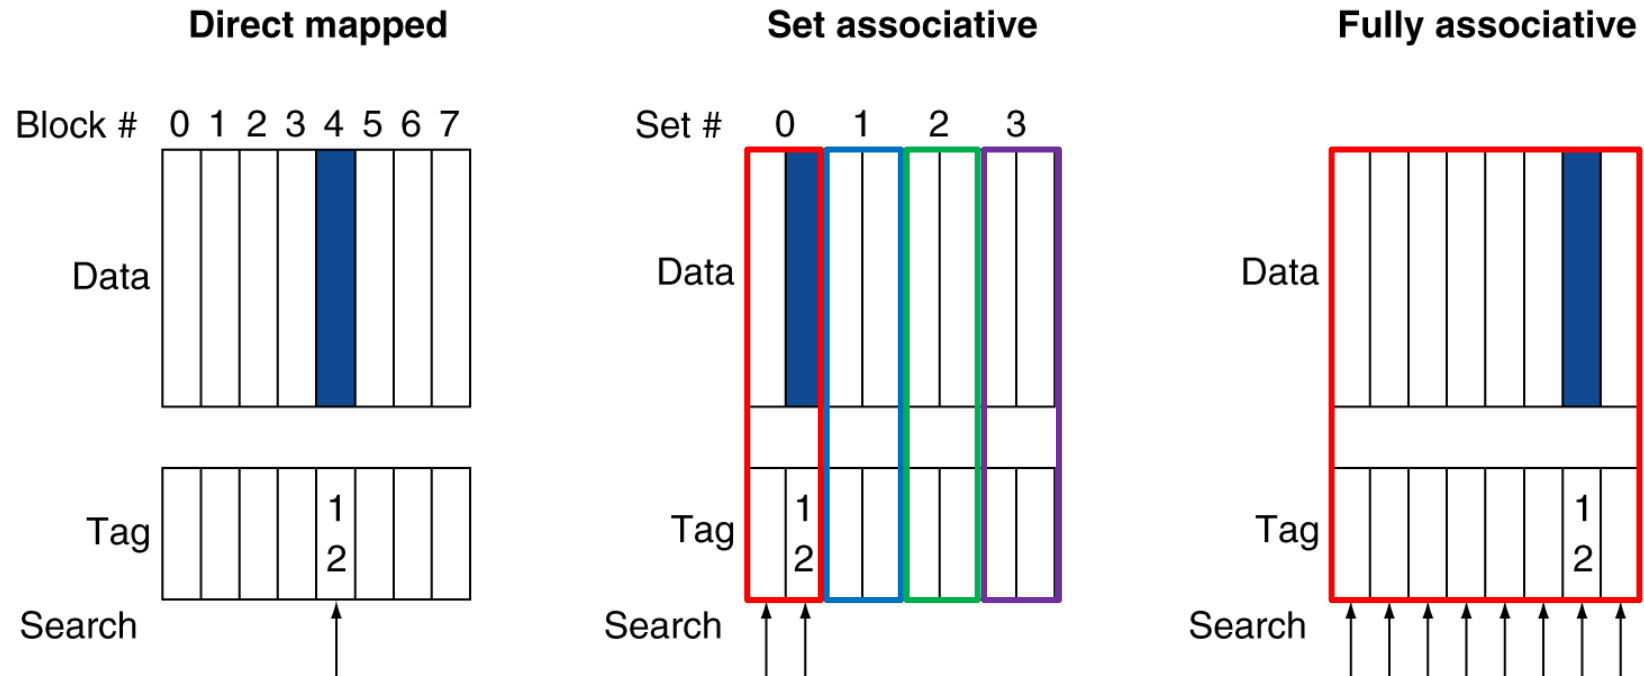
- ❖ Large cache has slow hit time

→ Solution: Use multiple levels of caches!

## ■ Software Optimization

- ❖ Solution: Strategically design data access pattern

# Associative Cache Example



# Associative Caches

---

## ■ Fully associative

- ❖ Allow a given block to be placed anywhere in the cache memory.
- ❖ Requires all locations to be searched at once (Otherwise, very slow..)
- ❖ A comparator for each entry (expensive)

## ■ *n*-way set associative

- ❖ Each **set** contains *n* locations (*n* blocks)
- ❖ Block number determines the set index
  - (Block number) modulo (#Sets in cache)
- ❖ Search all entries **in a given set** at once
- ❖ *n* comparators (less expensive)
- ❖ *n* may not be a power-of-two

# Spectrum of Associativity

- For a cache with 8 block locations

**One-way set associative  
(direct mapped)**

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

**Two-way set associative**

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

**Four-way set associative**

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

**Eight-way set associative (fully associative)**

Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

# Associativity Example

- Compare 4-block caches
  - ❖ Direct mapped / 2-way set associative / fully associative
  - ❖ Block address access sequence: 0, 8, 0, 6, 8
- Direct mapped

Block address	Cache index	Hit/miss	Cache content after access			
			0	1	2	3
0	0	miss	Mem[0]			

# Associativity Example

- Fully associative

Block address		Hit/miss	Cache content after access			
0		miss	Mem[0]			

# Associativity Example

- 2-way set associative

Block address	Cache index	Hit/miss	Cache content after access			
			Set 0		Set 1	
0	0	miss	Mem[0]			




# How Much Associativity

---

- Increased associativity *generally* decreases miss rate
  - ❖ But with diminishing returns
- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
  - ❖ 1-way: 10.3%
  - ❖ 2-way: 8.6%
  - ❖ 4-way: 8.3%
  - ❖ 8-way: 8.1%

# Set Associative Cache Organization



# Replacement Policy

---

- Direct mapped: no choice
- Set associative
  - ❖ Choose a non-valid (empty) location, if there is any.
  - ❖ Otherwise, need to choose one of the blocks in the set as the **victim**
- The best choice? The block that will be used in the most distant future.
  - ❖ Theoretically true, but impossible to know

# Replacement Policies for Set Associative 1

---

- Least-recently used (LRU)
  - ❖ Choose the one unused for the longest time
    - Simple for 2-way, manageable for 4-way, too hard beyond that
- Pseudo LRU (for 4+ ways)
  - ❖ Use an approximation
  - ❖ e.g., A tree-like structure that has an LRU bit per 2 entries.

# Replacement Policies for Set Associative 2

---

- Sequential (Round Robin)
  - ❖ Select the replacement victim in the circular order (i.e.,  $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 0 \rightarrow 1 \dots$ )
- Random
  - ❖ Similar performance as LRU if the associativity is high enough.

# Cache Example

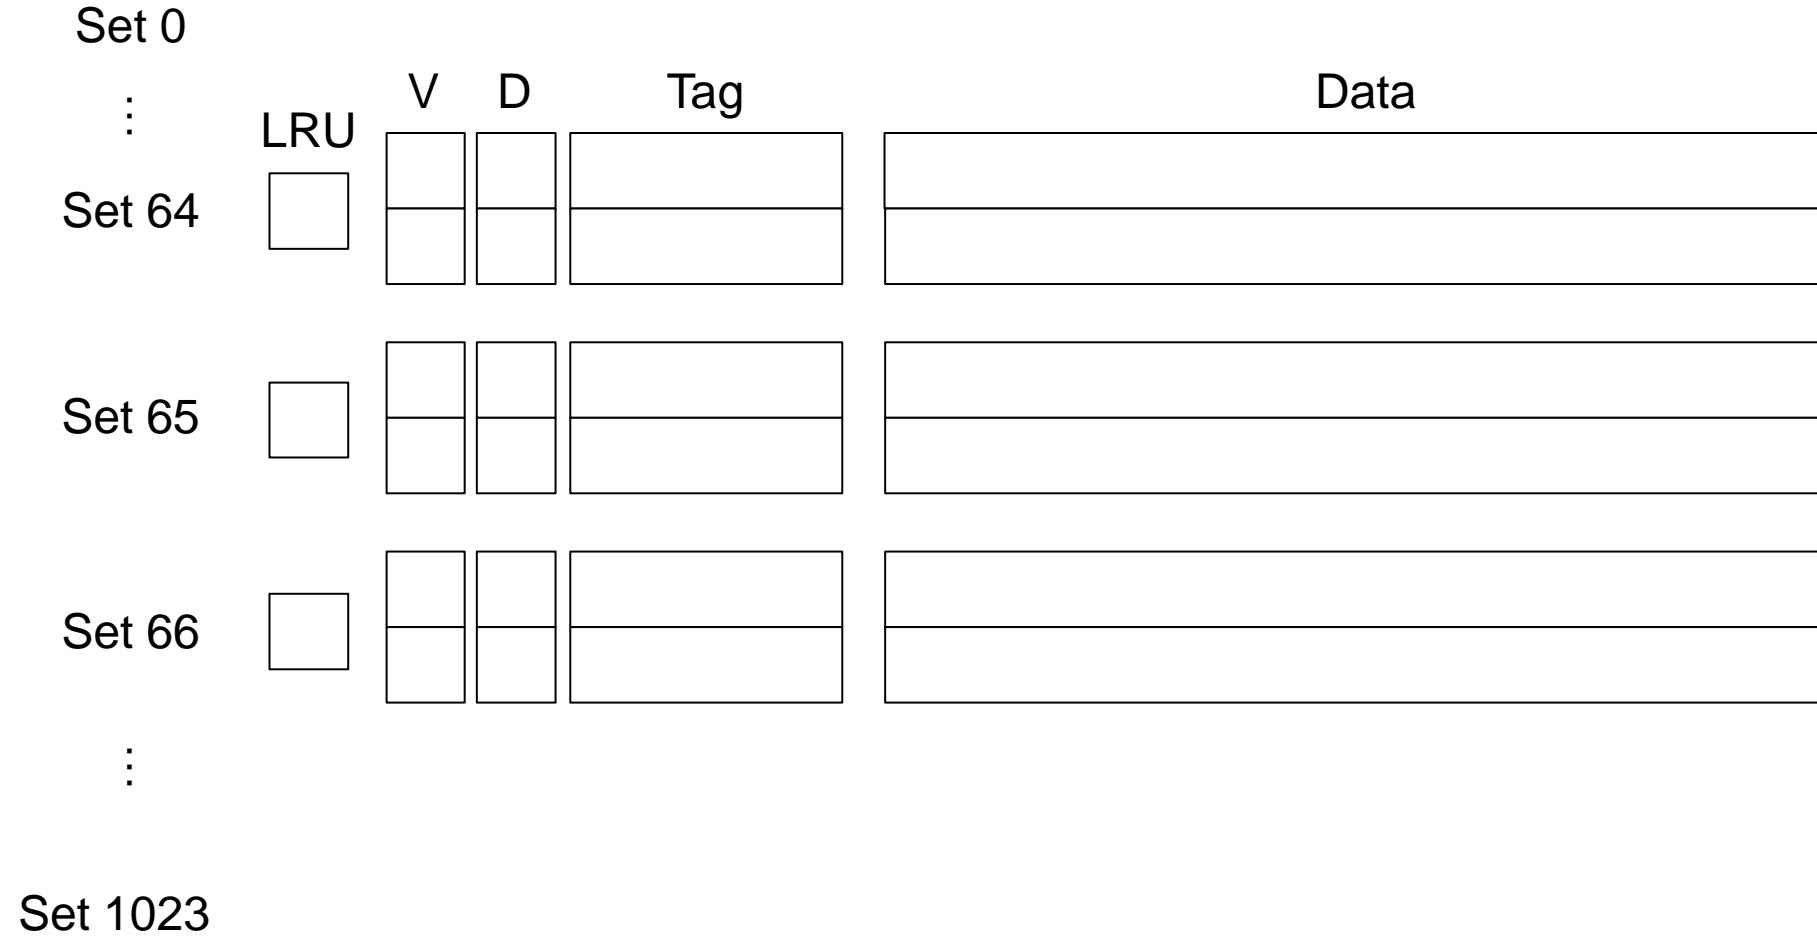
---

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

address

- Load 0x20001000
- Store 0x36801038
- Load 0x20001080
- Load 0x20001008
- Load 0x40811024
- Load 0x53841010
- Store 0x53841020
- Load 0x20001008


# Cache Example



# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>






# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.


	address	tag	index	offset	
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss



# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset	
■ Load	0x20001000	0x2000	$0001000000_2 \rightarrow 64_{10}$	$000000_2$	miss
■ Store	0x36801038	0x3680	$0001000000_2 \rightarrow 64_{10}$	$111000_2$	miss



# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.


	address	tag	index	offset	
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Store	0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■ Load	0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss

	LRU	V	D	Tag
Set 64	<div>↗</div>	1	0	0x2000
		1	1	0x3680
Set 66	<div>↘</div>	1	0	0x2000

# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

		address	tag	index	offset	
■	Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■	Store	0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■	Load	0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss
■	Load	0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	hit



# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset	
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Store	0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■ Load	0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Load	0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	hit
■ Load	0x40811024	0x4081	0001000000 <sub>2</sub> →64 <sub>10</sub>	100100 <sub>2</sub>	miss, write back 0x36801038

	LRU	V	D	Tag
Set 64	↗	1	0	0x2000
		1	0	0x4081
Set 66	↘	1	0	0x2000

# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset	
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Store	0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■ Load	0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Load	0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	hit
■ Load	0x40811024	0x4081	0001000000 <sub>2</sub> →64 <sub>10</sub>	100100 <sub>2</sub>	miss, write back 0x36801038
■ Load	0x53841010	0x5384	0001000000 <sub>2</sub> →64 <sub>10</sub>	010000 <sub>2</sub>	miss

	LRU	V	D	Tag
Set 64	<div> <div></div> <div>↘</div> </div>	1	0	0x5384
		1	0	0x4081
Set 66	<div> <div></div> <div>↘</div> </div>	1	0	0x2000

# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset	
■ Load	0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Store	0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■ Load	0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss
■ Load	0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	hit
■ Load	0x40811024	0x4081	0001000000 <sub>2</sub> →64 <sub>10</sub>	100100 <sub>2</sub>	miss, write back 0x36801038
■ Load	0x53841010	0x5384	0001000000 <sub>2</sub> →64 <sub>10</sub>	010000 <sub>2</sub>	miss
■ Store	0x53841020	0x5384	0001000000 <sub>2</sub> →64 <sub>10</sub>	100000 <sub>2</sub>	hit

	LRU	V	D	Tag
Set 64	<div> <div></div> <div>↘</div> </div>	1	1	0x5384
		1	0	0x4081
Set 66	<div> <div></div> <div>↘</div> </div>	1	0	0x2000

# Cache Example

- 128KB data capacity, 64 bytes per block, 32bit address
- 2-way set associative, LRU replacement
- Store policy: write-back, write-allocate.

	address	tag	index	offset	
■	Load 0x20001000	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	000000 <sub>2</sub>	miss
■	Store 0x36801038	0x3680	0001000000 <sub>2</sub> →64 <sub>10</sub>	111000 <sub>2</sub>	miss
■	Load 0x20001080	0x2000	0001000010 <sub>2</sub> →66 <sub>10</sub>	000000 <sub>2</sub>	miss
■	Load 0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	hit
■	Load 0x40811024	0x4081	0001000000 <sub>2</sub> →64 <sub>10</sub>	100100 <sub>2</sub>	miss, write back 0x36801038
■	Load 0x53841010	0x5384	0001000000 <sub>2</sub> →64 <sub>10</sub>	010000 <sub>2</sub>	miss
■	Store 0x53841020	0x5384	0001000000 <sub>2</sub> →64 <sub>10</sub>	100000 <sub>2</sub>	hit
■	Load 0x20001008	0x2000	0001000000 <sub>2</sub> →64 <sub>10</sub>	001000 <sub>2</sub>	miss

	LRU	V	D	Tag
Set 64	↗	1	1	0x5384
		1	0	0x2000
Set 66	↘	1	0	0x2000



# Multi-level Caches

---

- Primary cache attached to CPU
  - ❖ Small, but fast
- Level-2 cache handles misses from the L-1 cache
  - ❖ Larger, slower, but still faster than the main memory
- Main memory handles L-2 cache misses
- Some high-end systems include L-3 cache
  - ❖ Sometimes called the last-level cache or LLC.

# Multi-level Cache Example

---

- Given
  - ❖ CPU base CPI = 1, clock rate = 4GHz
  - ❖ Miss rate/instruction = 2%
  - ❖ Main memory access time = 100ns
  
- With just primary cache
  - ❖ Miss penalty =  $100\text{ns} / 0.25\text{ns} = 400$  cycles
  - ❖ Effective CPI =  $1 \times 100\% + 400 \times 2\% = 9$

# Example (cont.)

---

- Now add L-2 cache
  - ❖ Access time = 5ns
  - ❖ **Global** miss rate to main memory = 0.5%
    - **Local** L2 cache miss rate?
- Primary miss with L-2 hit
  - ❖ Penalty =  $5\text{ns}/0.25\text{ns} = 20$  cycles
- Primary miss with L-2 miss
  - ❖ **Extra** penalty = 400 cycles
- $\text{CPI} = 1 \times 100\% + 20 \times 2\% + 400 \times 0.5\% = 3.4$
- Performance ratio =  $9/3.4 = 2.6\times$  faster with L2

# Multi-level Cache Considerations

---

- Primary (L-1) cache
  - ❖ Focus on minimal hit time
- L-2 cache
  - ❖ Focus on low miss rate to avoid main memory access
  - ❖ Hit time has less overall impact


# Sources of Misses – The Three C's

---

- **Compulsory** misses (a.k.a. cold start misses)
  - ❖ First access to a block
- **Capacity** misses
  - ❖ Due to finite cache size
  - ❖ A replaced block is later accessed again
- **Conflict** misses (a.k.a. collision misses)
  - ❖ In a non-fully associative cache
  - ❖ Due to competition for entries in a set
  - ❖ Do not occur in fully associative caches

# VIRTUAL MEMORY

# Virtual Memory



# Virtual Address

---

User application

load  
store

address 0x4124

← **Virtual address**

**Translate**

address 0x30124

← **Physical address**

Memory (cache)



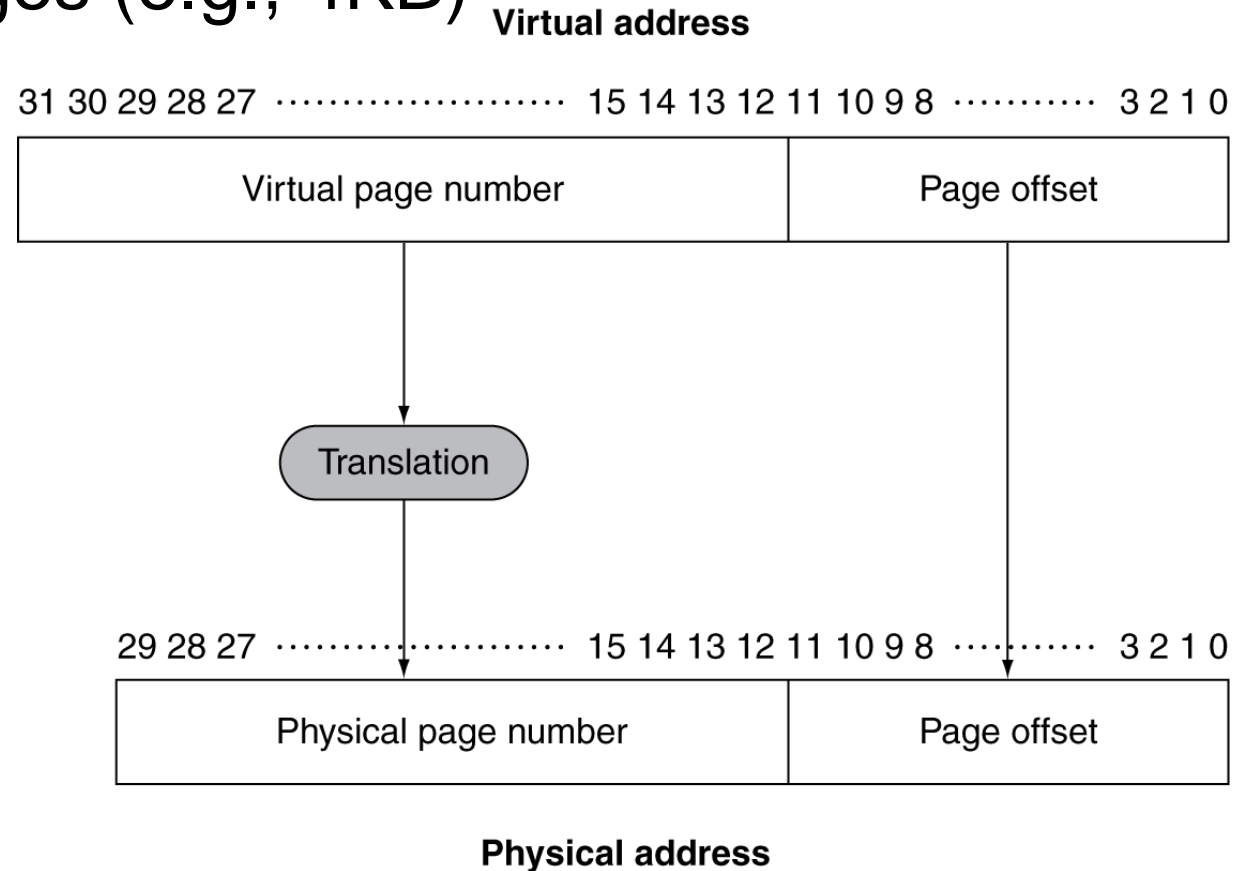
# Why Virtual Memory?

---

- Each program gets a private virtual address space
  - ❖ Multiple programs share main memory
  - ❖ Do not need to be programmed for a specific memory size
- If there isn't enough main memory, relocates some data to storage (disk).
  - ❖ Managed jointly by CPU hardware and the operating system (OS)

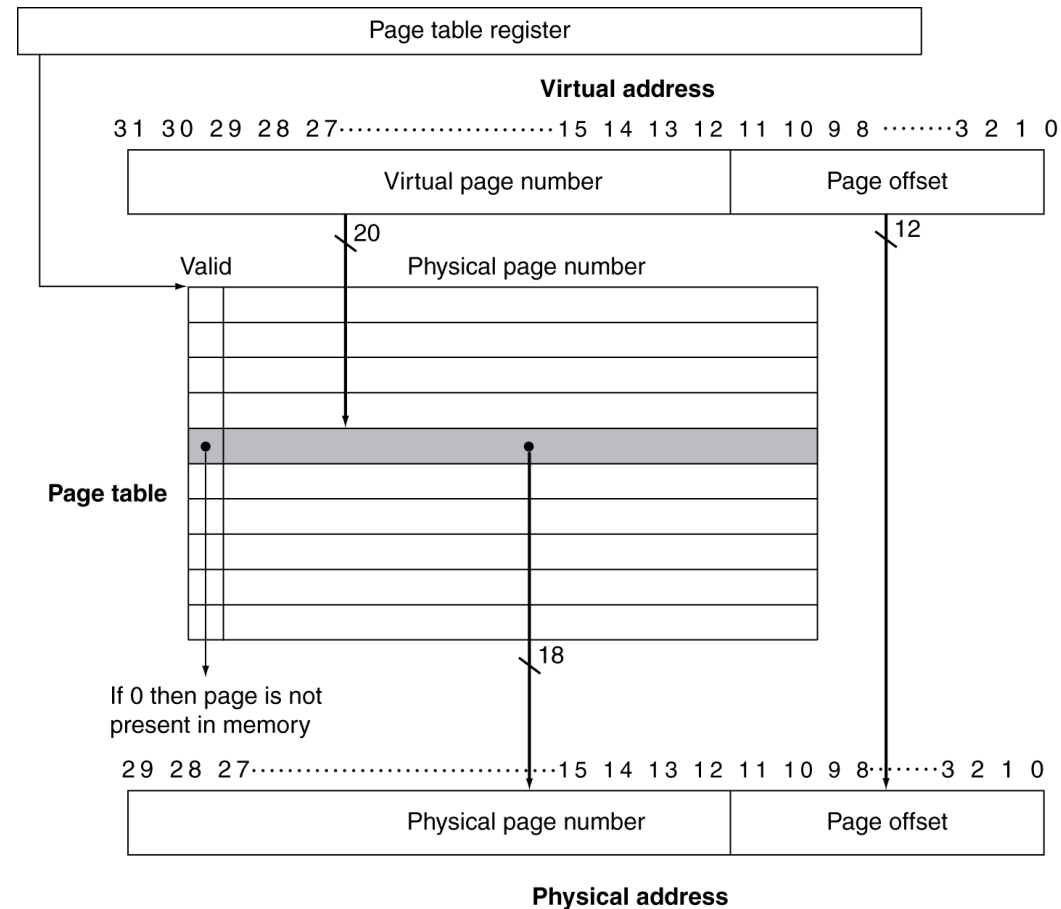
# Address Translation

- Fixed-size pages (e.g., 4KB)



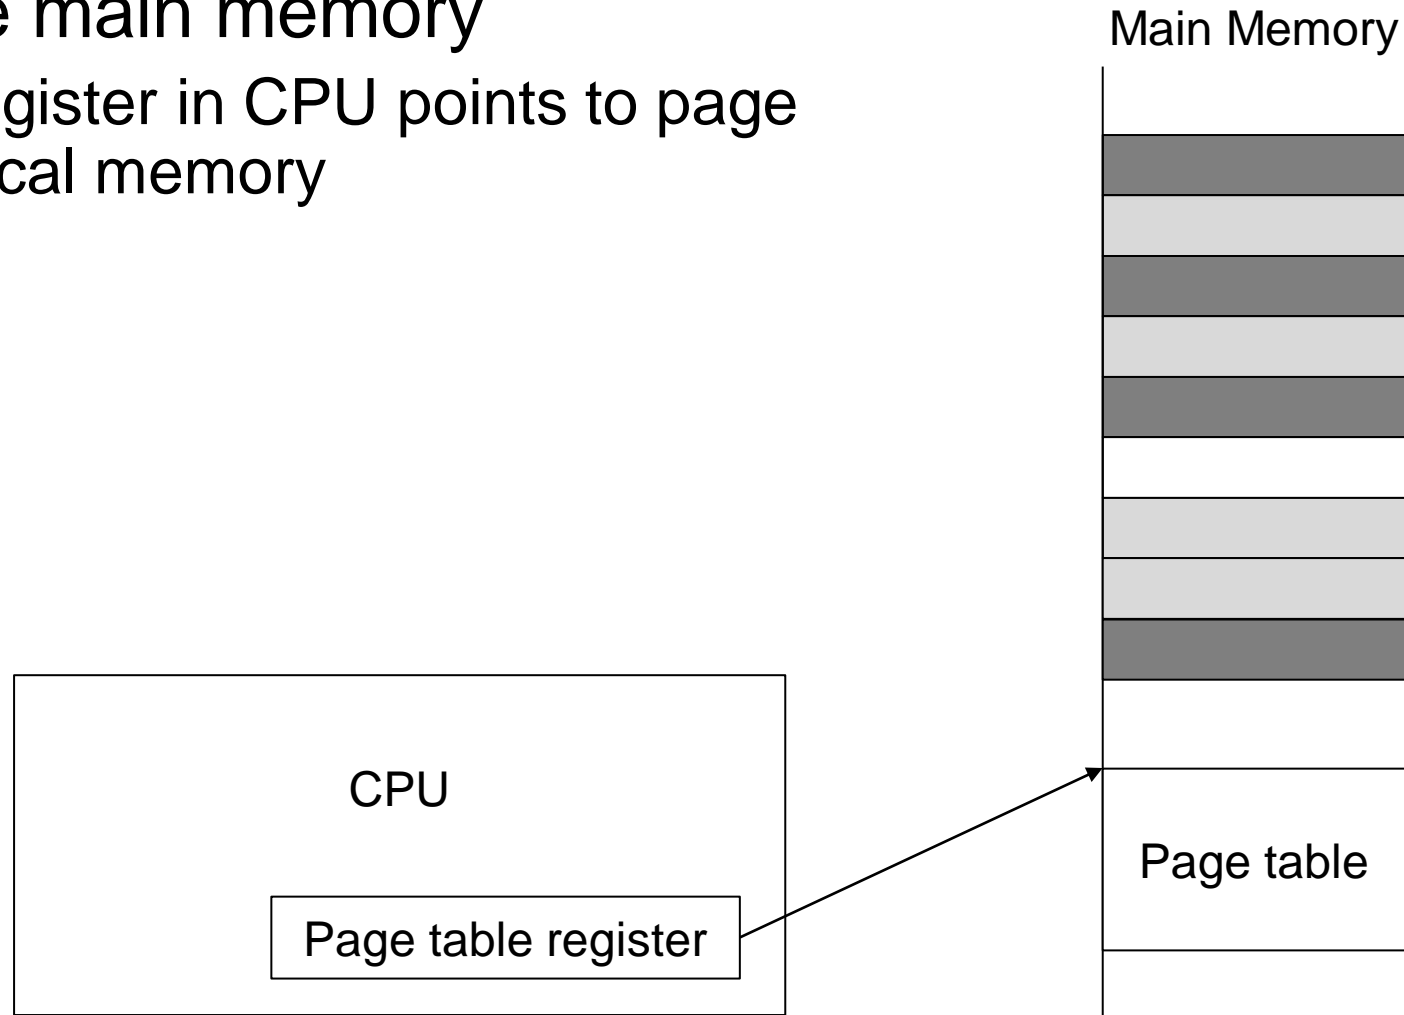
# Translation Using a Page Table

- Stores page mapping information
  - ❖ Array of **page table entries**
  - ❖ Indexed by the virtual page number



# Page Table in Memory

- Resides in the main memory
  - ❖ Page table register in CPU points to page table in physical memory



# Page Table Entry (PTE)

- If page is present in memory
  - ❖ PTE stores the physical page number
  - ❖ Plus, other status bits (valid, dirty, referenced, ...)
- If page is not present (i.e., valid bit is 0)
  - ❖ This means the page is in the disk or not mapped to memory .
  - ❖ Accessing the page → **Page fault**

## Page table

PTE for virtual page 0

PTE for virtual page 1

PTE for virtual page 2

⋮

⋮

⋮

4 bytes			
V	D	R	Physical page number
V	D	R	Physical page number
V	D	R	Physical page number

# Page Fault Handler

---

- **Page fault** is handled by OS
  - ❖ Locate page on disk
  - ❖ Choose a page to replace
    - If dirty, write to disk first
  - ❖ Read page into memory and update page table

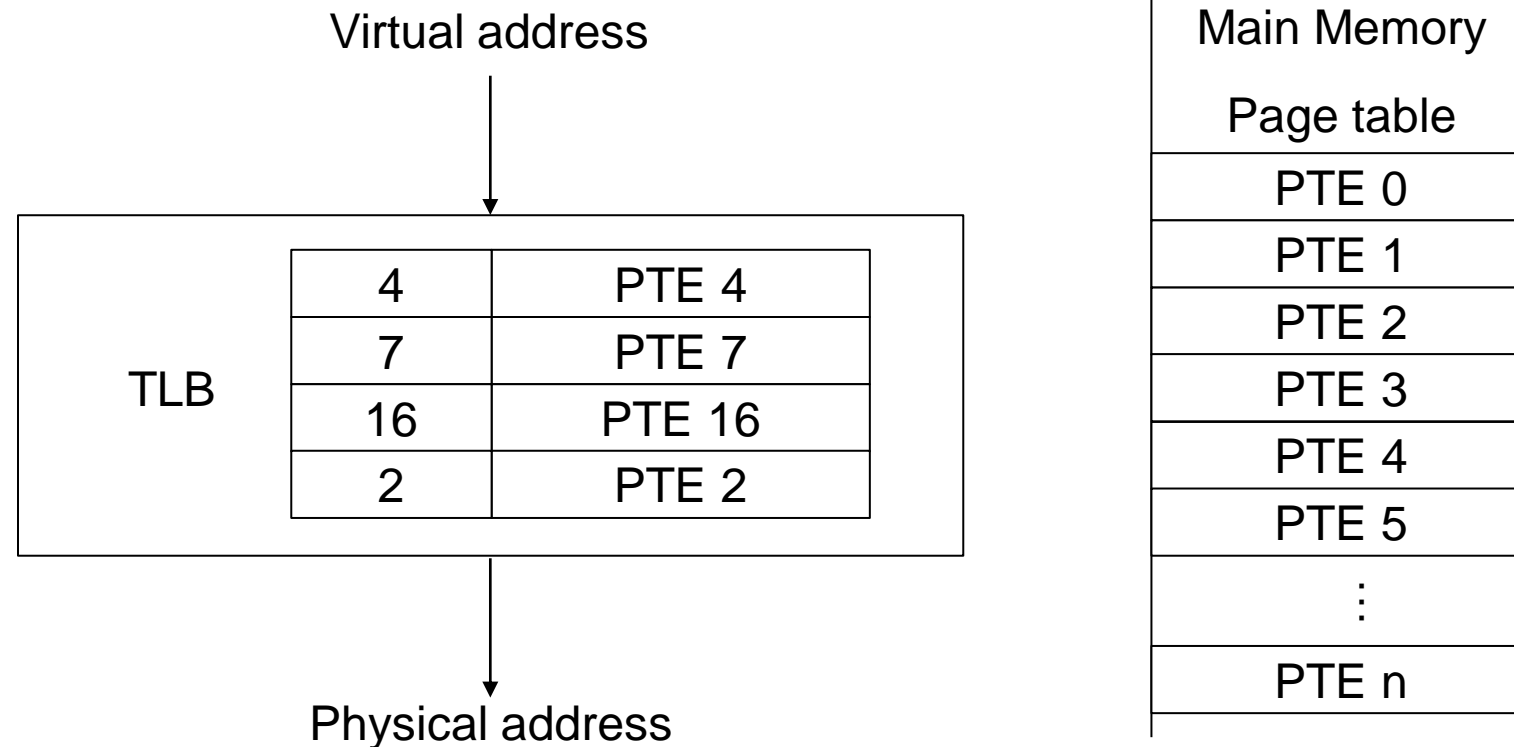
# Fast Translation Using a TLB

---

- Address translation would appear to require extra memory references
  - ❖ One to access the PTE
  - ❖ Then the actual memory access
- Usually, accesses to page table has good locality patterns
  - ❖ Use a fast cache of PTEs within the CPU
  - ❖ Called a Translation Look-aside Buffer (TLB)

# TLB

- Store frequently accessed PTEs in a faster memory
- Typical: 16–512 PTEs
- Fully associative, n-way set associative, ...





# TLB Misses

---

- TLB miss: The page table entry (PTE) is not in TLB
  - ❖ This can be handled by the CPU (hardware TLB management) or by the OS (software TLB management)
    - x86: hardware TLB management
    - MIPS: software TLB management
- If PTE says...
  - ❖ The page is in memory
    - Load the PTE from memory to TLB and retry the memory access
  - ❖ The page is not in memory (page fault)
    - OS handles fetching the page and updating the page table

# Memory Access Scenarios (Data load)

